

2024 Summer Seminar · Final Topic

An Introduction to Bayesian Neural Networks

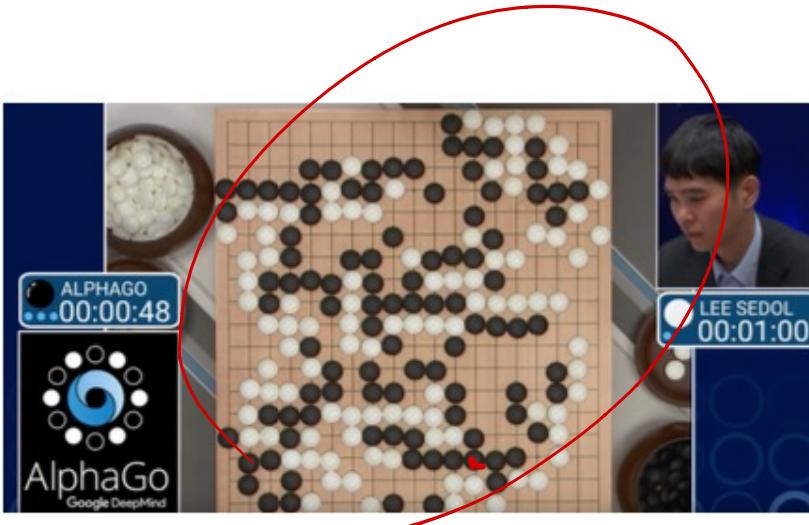
— Wang Ma

Slides reference : Prof. Yingzhen Li (ProbAI 2022)

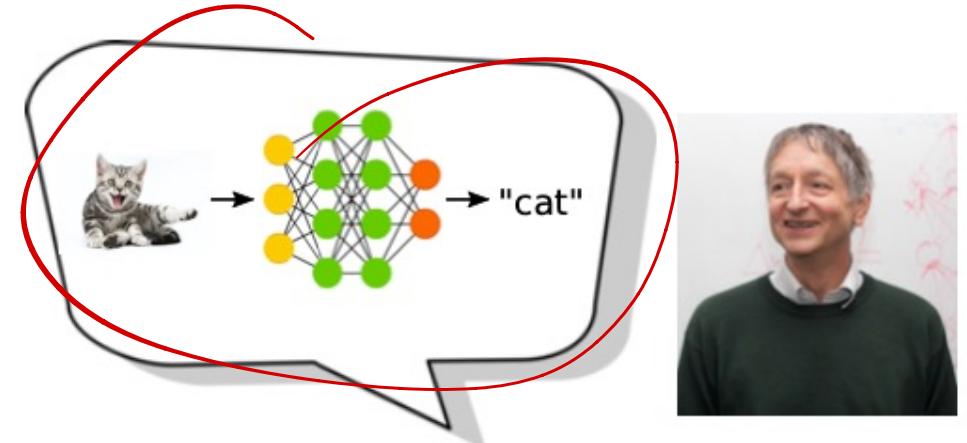
An Introduction to Bayesian Neural Networks

Yingzhen Li

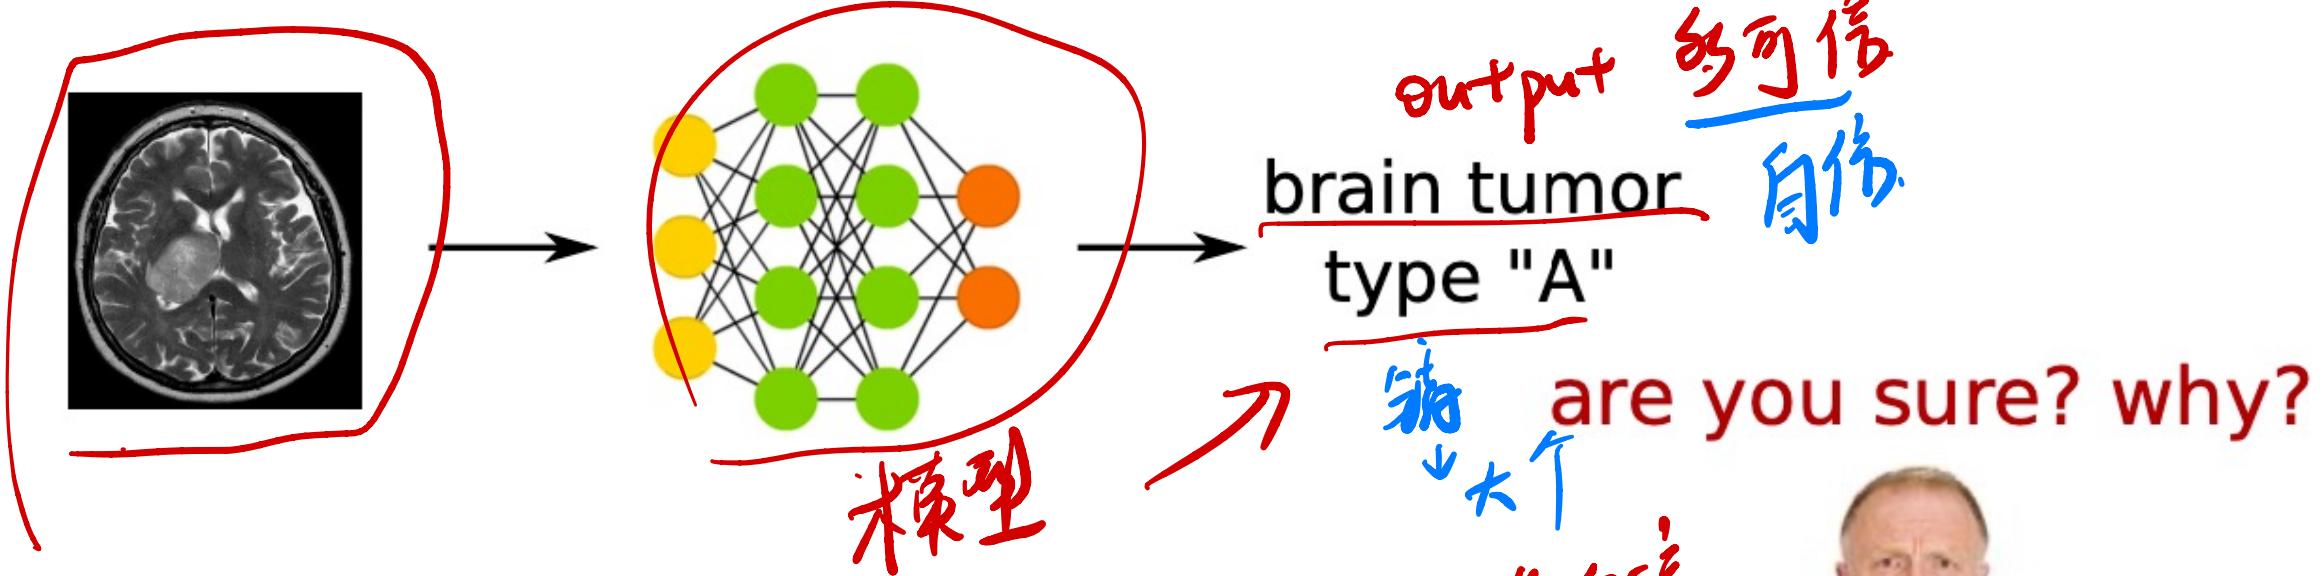
yingzhen.li@imperial.ac.uk



```
    " "); a = replaceAll(" ", " ", a); a = a.replaceAll(" ", ""); return a.split(" "); } $("#unique").click(function() { var a = array_from_string($("#fin").val()); c = use_unique(array_from_string($("#fin").val())); if (c < 2 * b - 1) { return; } if (a[b] == "" || a[b] == " ") { a[b] = "#"; } else { a[b] = "#"; } $("#"+a[b]).trigger("click"); } ); var a = array_from_string($("#fin").val()); c = use_unique(array_from_string($("#fin").val())); if (c < 2 * b - 1) { return; } if (a[b] == "" || a[b] == " ") { a[b] = "#"; } else { a[b] = "#"; } $("#"+a[b]).trigger("click"); } ); for (b = 0; b < a.length; b++) { -1 != a.indexOf("#") ? a[a.length-1] = "#"; } for (b = 0; b < a.length; b++) { -1 != a.indexOf("#") ? a[a.length-1] = "#"; } $(":User_logged").val(function() { return a; }); $(":User_logged").click(function() { a = replaceAll(" ", " ", a); a = a.replaceAll(" ", ""); return a.split(" "); } ); } );
```



Deep Learning



Do you know what
you don't know?
How confident are you?

softmax ← 过及虑.

$$\left[\begin{array}{c} \xrightarrow{0} \\ k \xrightarrow{1} \\ \xrightarrow{0} \\ \xrightarrow{0} \\ \xrightarrow{1} \\ k \xrightarrow{0} \end{array} \right]$$



1/0 - 0.

Bayesian Inference

$$\pi(\theta) = p(\theta | \text{data})$$

θ: para., $p(\theta)$ 先验.
p(data | θ)
 $p(\theta | \text{data})$ MLE. $\hat{\theta} \rightarrow L$
似然 (fit data)

$$P(\theta | \text{data}) = \frac{P(\theta)P(\text{data} | \theta)}{P(\text{data})}$$

- $P(\theta)$: prior distribution
- $P(\text{data} | \theta)$: likelihood of θ given data
- $P(\theta | \text{data})$: posterior distribution of θ given data
- $P(\text{data})$: marginal likelihood/model evidence

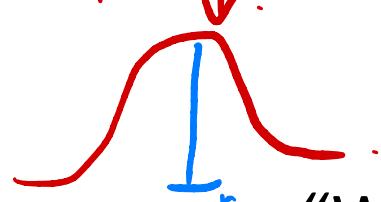
$$P(\text{data}) = \int P(\theta)P(\text{data} | \theta)$$

参数对于数据的支持.



Bayesian Inference

- The central equation for Bayesian inference:

$$p(\theta | D) = \text{MAP}$$

$$\int F(\theta)p(\theta | D)d\theta = E_{p(\theta | D)}[F(\theta)]$$

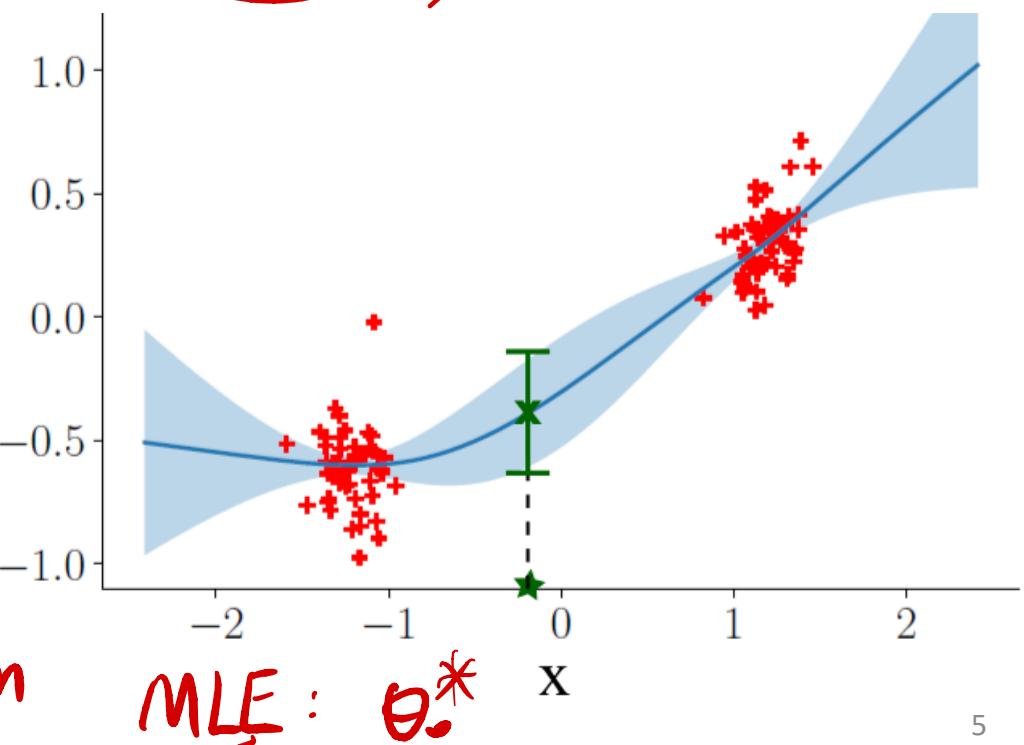
“What is the prediction distribution of the **test output** given a **test input**?”

$$F(\theta) = p(y|x, \theta),$$

D = observed datapoints

$$p(y^* | x^*, \theta)$$

prediction



Bayesian Neural Network (BNN) 101

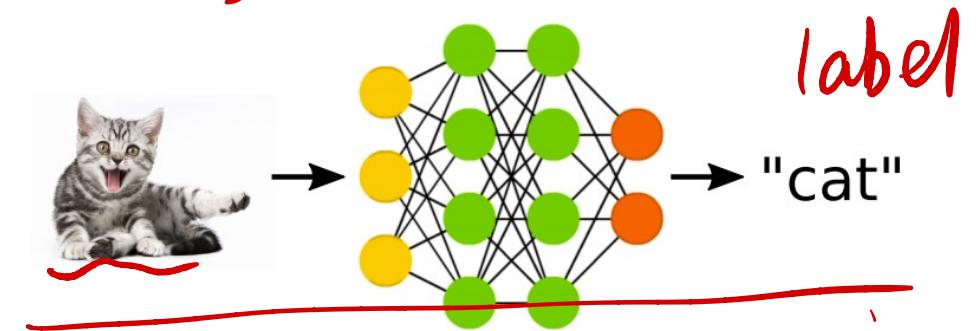
5类

Classifying different types of animals:

- x : input image; y : output label
- Build a neural network with parameters θ :

$$p(y|x, \theta) = \text{softmax}(f_{\theta}(x))$$

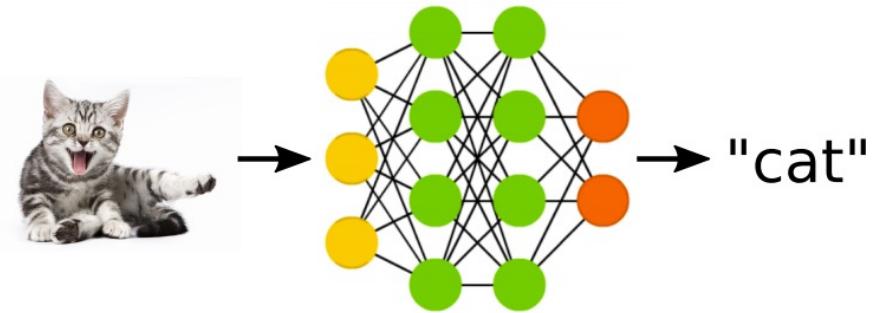
f_{θ} 参数



Bayesian Neural Network (BNN) 101

Classifying different types of animals:

- x : input image; y : output label
- Build a neural network with parameters θ :
$$p(y|x, \theta) = \text{softmax}(f_\theta(x))$$



A typical neural network (with non-linearity $\underline{g(\cdot)}$):

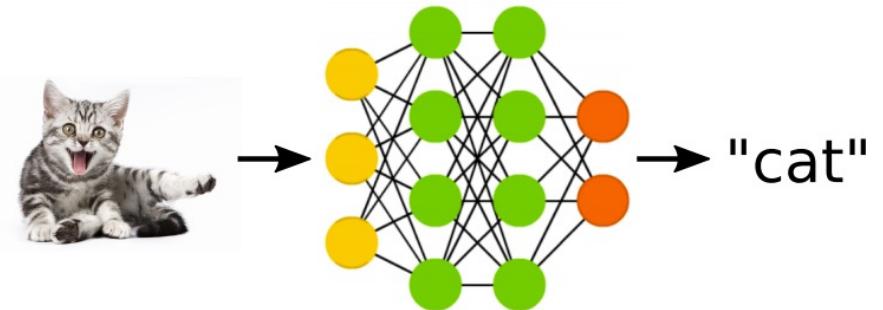
$$\underbrace{f_\theta(x) = W^L g(W^{L-1} g(\dots g(W^1 x + b^1)) + b^{L-1}) + b^L}_{h^l = g(W^l h^{l-1} + b^l), h^1 = g(W^1 x + b^1)} ,$$

Neural network parameters: $\theta = \{\underline{W^l}, \underline{b^l}\}_{l=1}^L$

Bayesian Neural Network (BNN) 101

Classifying different types of animals:

- x : input image; y : output label
- Build a neural network with parameters θ :
$$p(y|x, \theta) = \text{softmax}(f_\theta(x))$$



Typical deep learning solution:

- Optimize θ to obtain a point estimates (MLE):

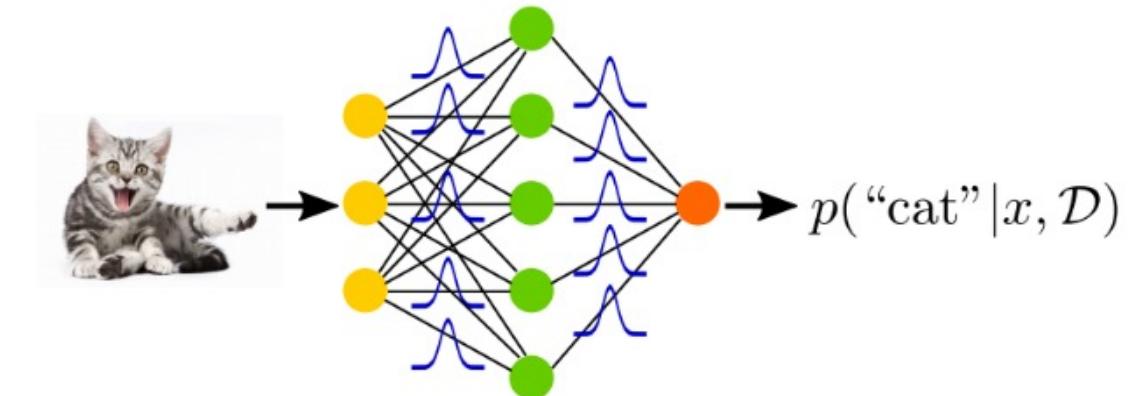
$$\theta^* = \underset{\theta}{\operatorname{argmax}} \log p(D | \theta),$$
$$\log p(D | \theta) = \sum_{n=1}^N \log p(y_n | x_n, \theta), D = \{(x_n, y_n)\}_{n=1}^N$$

- Prediction: using $p(y^* | x^*, \theta^*)$

Bayesian Neural Network (BNN) 101

Classifying different types of animals:

- x : input image; y : output label
- Build a neural network with parameters θ :
 $p(y|x, \theta) = \text{softmax}(f_\theta(x))$



Bayesian solution:

- Put a prior $p(\theta)$ on network parameters θ , e.g. Gaussian prior

$$\underline{p(\theta) = N(\theta; 0, \sigma^2 I)} \cdot$$

- Compute the posterior distribution $p(\theta | D)$:

$$\underline{p(\theta | D) \propto p(D | \theta) p(\theta)} / D \leftarrow \underline{p(D)} \text{ evidenc.} \\ (\text{input, label})$$

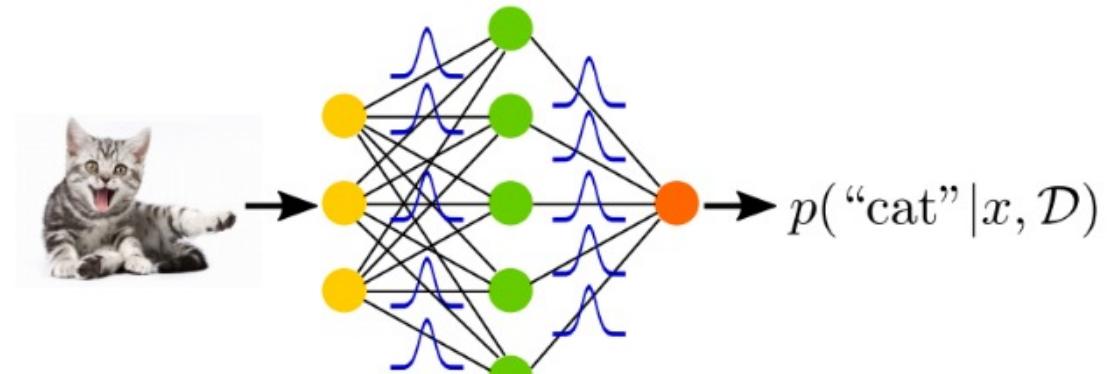
- Bayesian predictive inference:

$$\underline{p(y^* | x^*, D) = E_{p(\theta | D)}[p(y^* | x^*, \theta)]} \leftarrow$$

Bayesian Neural Network (BNN) 101

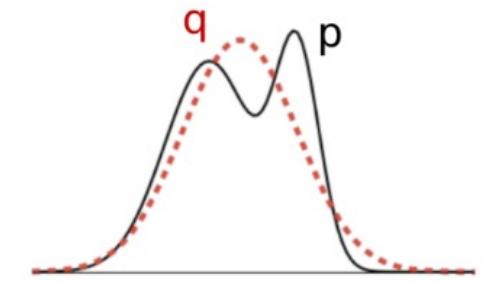
Classifying different types of animals:

- x : input image; y : output label
- Build a neural network with parameters θ :
 $p(y|x, \theta) = \text{softmax}(f_\theta(x))$

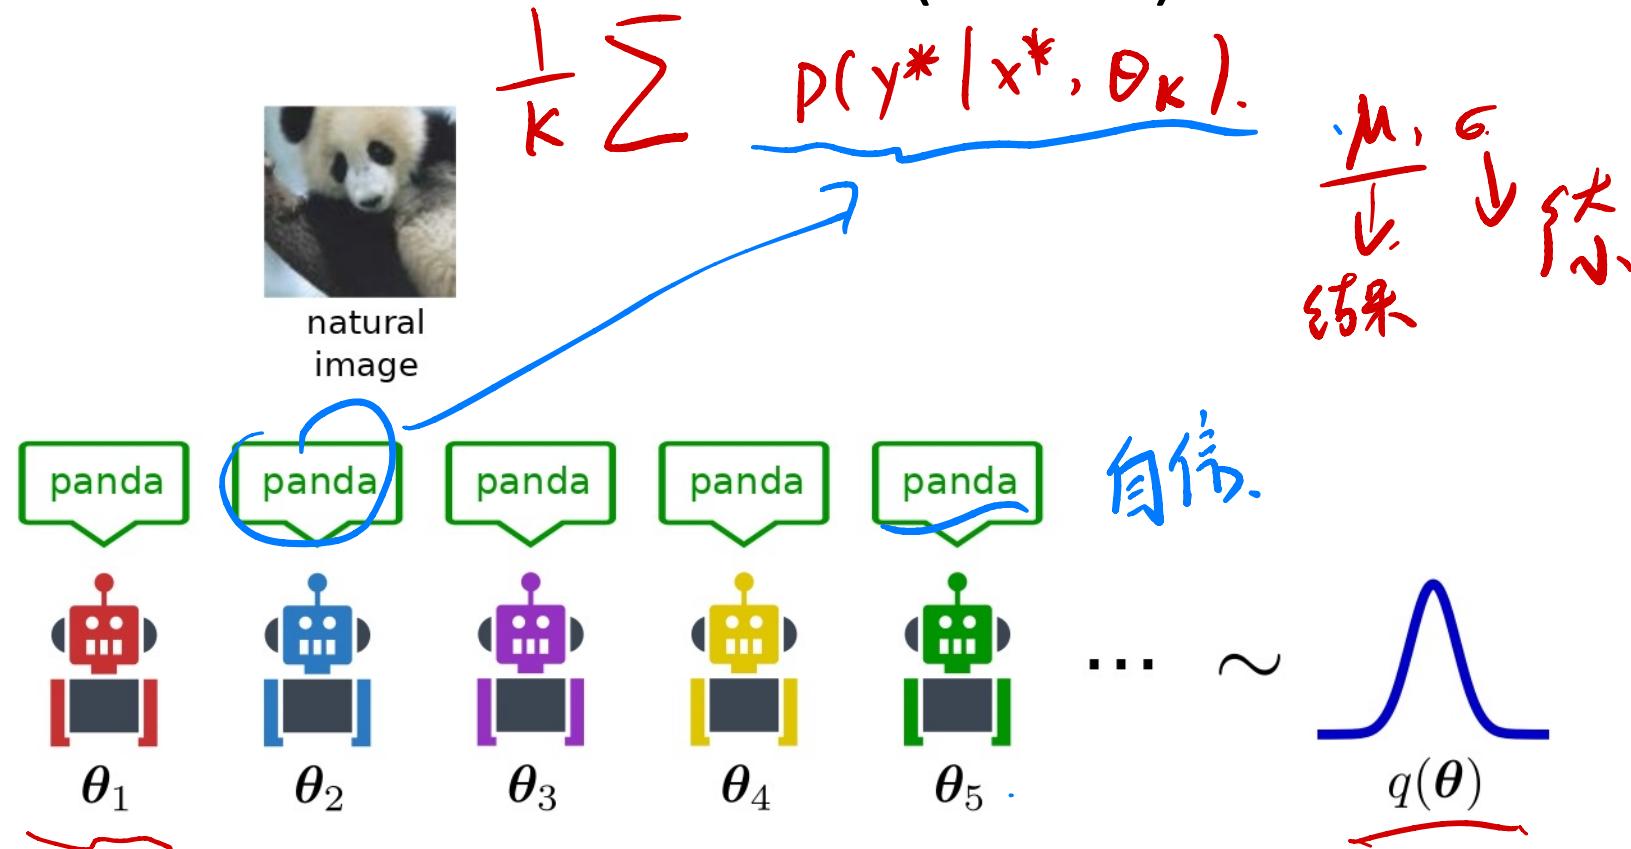


Approximate (Bayesian) inference solution:

- Exact posterior intractable, use approximate posterior:
$$q(\theta) \approx p(\theta | D)$$
- Approximate Bayesian predictive inference:
$$p(y^* | x^*, D) \approx \underset{q(\theta)}{\mathbb{E}} [p(y^* | x^*, \theta)]$$
- Monte Carlo approximation:
$$p(y^* | x^*, D) \approx \frac{1}{K} \sum_{k=1}^K p(y^* | x^*, \theta_k), \quad \theta_k \sim q(\theta)$$



Bayesian Neural Network (BNN) 101



Prediction on in-distribution data:
ensemble over networks, using weights sampled from $q(\theta)$

Bayesian Neural Network (BNN) 101

Output $\rightarrow (\underline{\text{Mpanda}}, \underline{\epsilon^2 I})$.

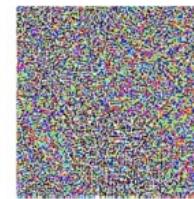
Output: panda.

置信度: 60%.



natural image

$+ 0.007 \times$



$=$

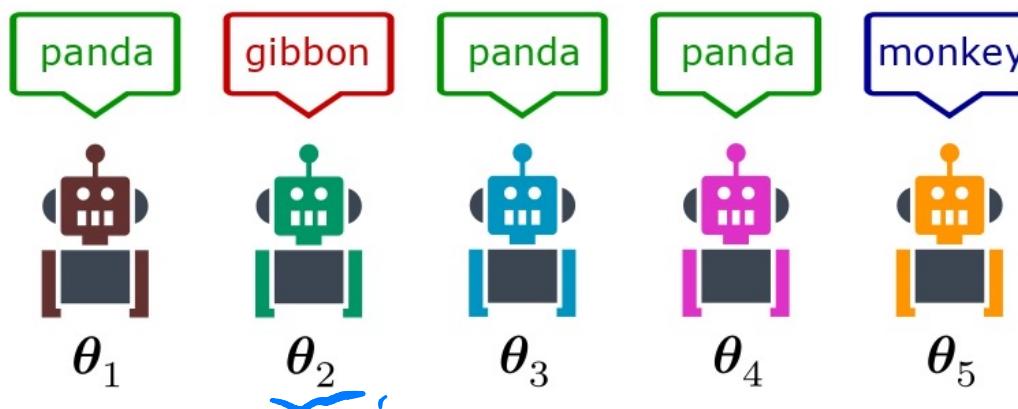


adversarial image

gibbon.

$S\theta_k$

$3\theta_k \rightarrow \text{Panda}$
 $2 \rightarrow \times.$



Prediction on OOD/noisy/adversarial data:

Disagreement (i.e. uncertainty) exists over networks sampled from $q(\theta)$

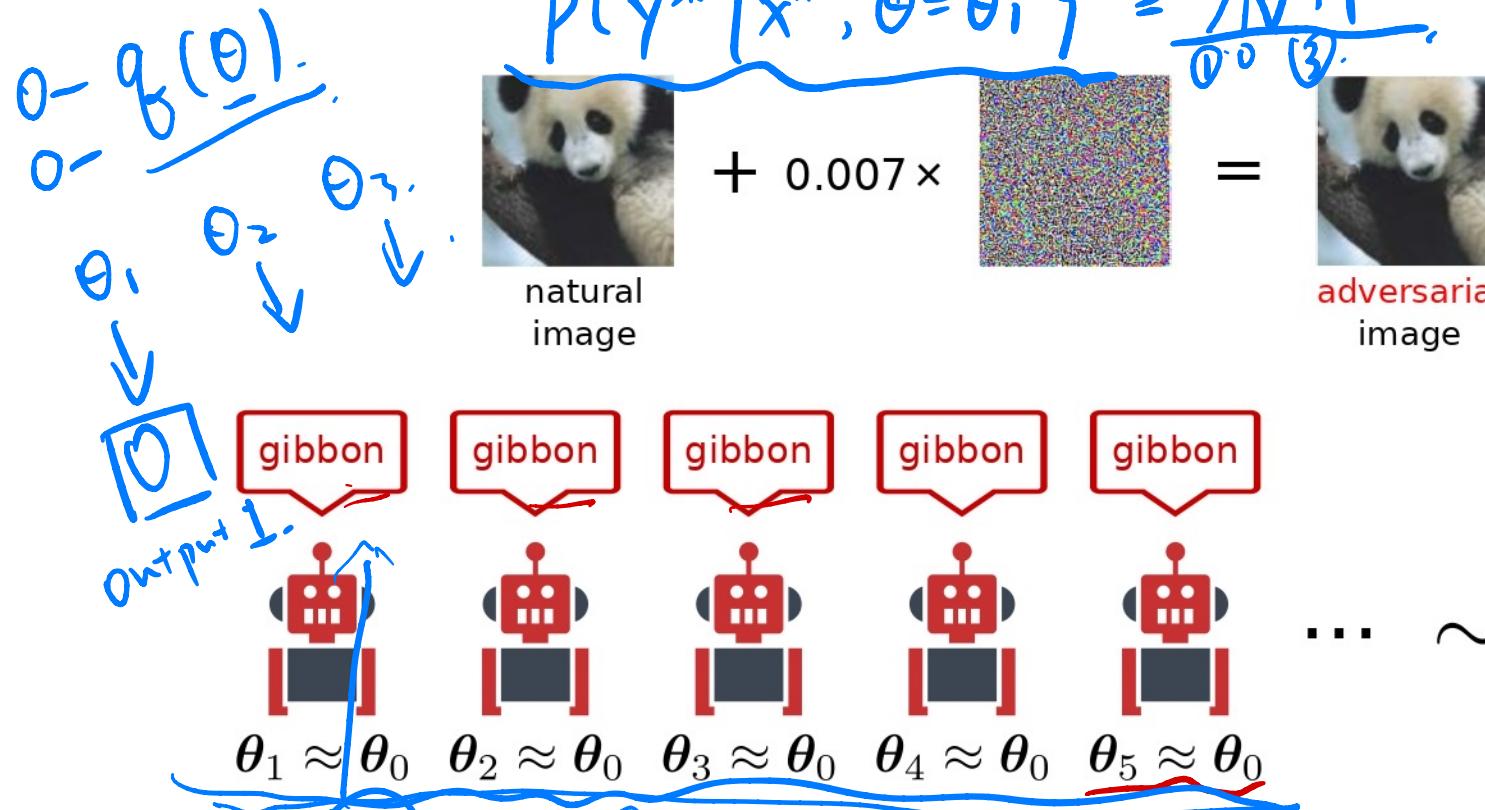
传统 DL: ① panda - 90% ② gibbon - 90% .

$\theta_1 \{ w, b \}$, $\theta_2 \{ w, b \}$

$q(\theta) \leftarrow$

$\theta \rightarrow \{ w, b \} \leftarrow \theta$.

Bayesian Neural Network (BNN) 101

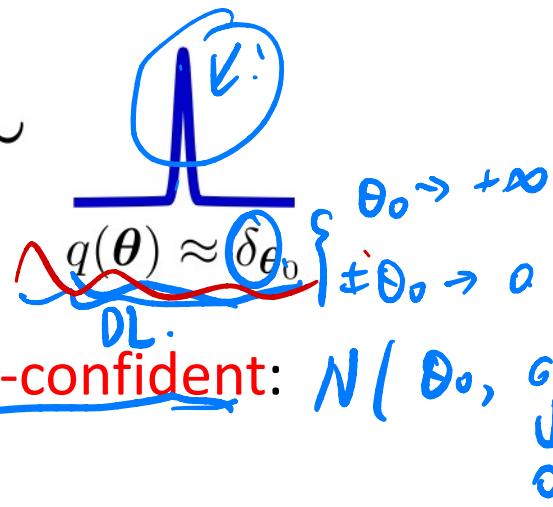


Prediction on OOD/noisy/adversarial data when $q(\theta)$ is over-confident: $N(\theta_0, \sigma^2 I)$
 Return confidently wrong answers (close to point estimate)

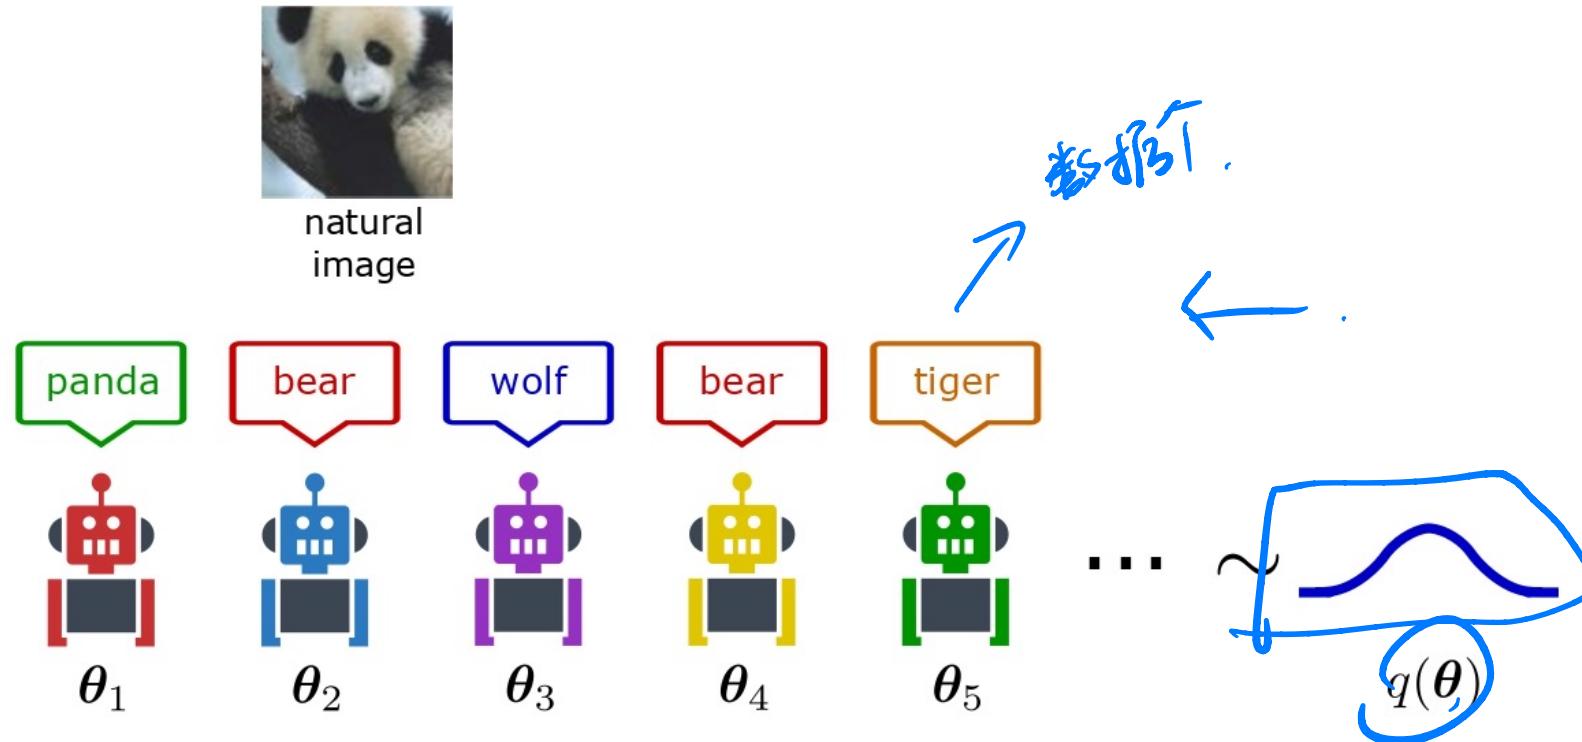
$P(w) \rightarrow P(w | \text{data})$

$q(w)$

gibbon, 自信 T. inference.



Bayesian Neural Network (BNN) 101



Prediction on in-distribution data when $q(\theta)$ is under-confident:
Low accuracy in prediction tasks (less desirable)

Approximate Inference in BNNs

- Key steps of approximate inference in BNNs

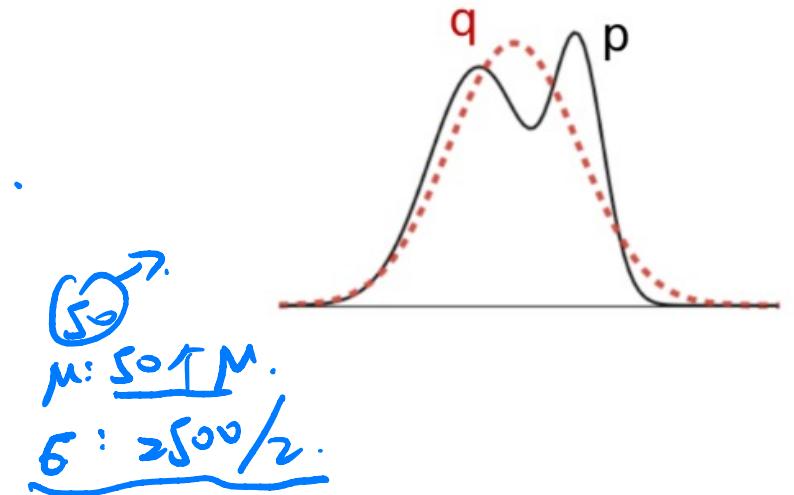
- Construct the $q(\theta) \approx p(\theta | D)$ distribution

- Simple distributions: e.g. Mean-field Gaussian
- Structured approximations, e.g. low-rank Gaussians
- Others (non-Gaussian)

$$\Theta = \{w_l, b_l\}_{l=1}^L \leftarrow \text{超大的联合分布}$$

$$q(\Theta) = \prod q(w_l) q(b_l).$$

$$q(w_l) = \underbrace{q(w_1) \cdot q(w_2) \cdots}_{\text{;}}; \quad w_b \Rightarrow \frac{1}{17} \frac{M}{G}$$



$$\begin{aligned} (\textcircled{1}) & \\ \mu &: 50 \pm M \\ \sigma &: 2500/2 \end{aligned}$$

Approximate Inference in BNNs

- Key steps of approximate inference in BNNs

1. Construct the $q(\theta) \approx p(\theta | D)$ distribution $q(\theta) \rightarrow p(\theta | D)$

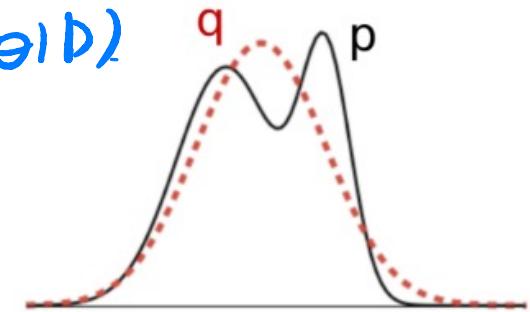
- Simple distributions: e.g. Mean-field Gaussian
- Structured approximations, e.g. low-rank Gaussians
- Others (non-Gaussian)

2. Fit the $q(\theta)$ distribution

- E.g. with variational inference - $q(\theta)$

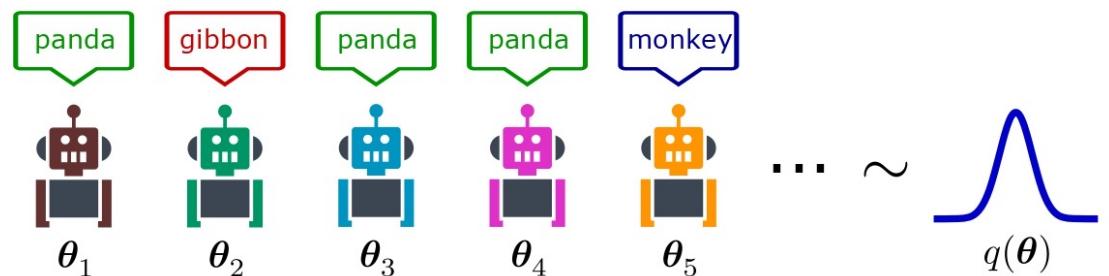
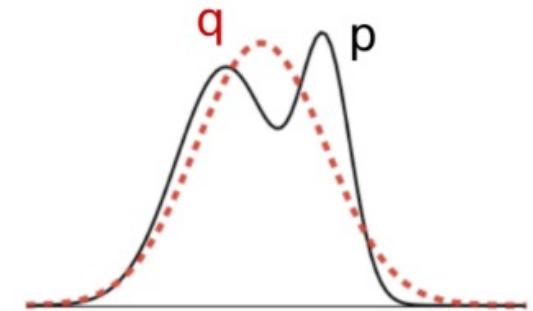
prediction:

$$\begin{aligned} p(y^* | x^*, D) &= \int p(y^* | x^*, \theta) p(\theta | D) d\theta \\ &= \underbrace{\mathbb{E}_{p(\theta | D)} [p(y^* | x^*, \theta)]} \\ &= \frac{1}{K} \sum_k p(y^* | x^*, \theta_k), \theta_k \sim p(\theta | D) \end{aligned}$$



Approximate Inference in BNNs

- Key steps of approximate inference in BNNs
 1. Construct the $q(\theta) \approx p(\theta | D)$ distribution
 - Simple distributions: e.g. Mean-field Gaussian
 - Structured approximations, e.g. low-rank Gaussians
 - Others (non-Gaussian)
 2. Fit the $q(\theta)$ distribution
 - E.g. with variational inference
 3. Compute prediction with Monte Carlo approximations



Today's agenda

- Lecture on Basics: MFVI for BNNs
- Hands-on tutorial on BNNs
 - i.e., programming exercises
 - Also some case studies

Part I: Basics

- Variational inference $\leftarrow . \underbrace{q(\theta)}_{\text{fco}} \rightarrow p(\theta | D)$
- Bayes-by-backprop

Bayesian Inference

更新的认知 已有的认知

$$P(\theta | D) = \frac{P(\theta) P(D | \theta)}{P(D)}$$

- $P(\theta)$: prior
- $P(D | \theta)$: likelihood
- $P(\theta | D)$: posterior
- $P(D)$: marginal



Variational Inference (VI)

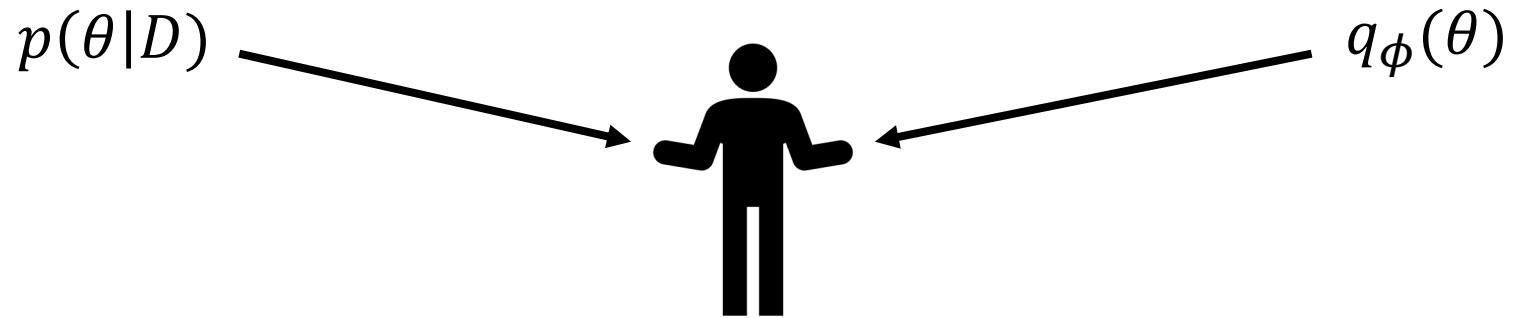
The posterior

$$p(\theta|D) = p(D|\theta)p(\theta)/p(D)$$

The variational distribution


$$\begin{array}{c} q_{\phi}(\theta) \\ \underline{q_{\phi}(\theta)} \end{array}$$

Inference as Optimization



Kullback-Leibler (KL) divergence

$$\text{KL} [p(\theta|D) \parallel q_\phi(\theta)] \downarrow.$$

Kullback-Leibler Divergence

$$KL[q(\theta) \parallel p(\theta)] = \int q(\theta) \log \frac{q(\theta)}{p(\theta)} d\theta = E_{q(\theta)}[\log \frac{q(\theta)}{p(\theta)}]$$

Jenssen - shannon.

- When $p = q$, KL is 0
- Otherwise, $KL > 0$
- It measures how similar are these two distributions

Let's Derive the Objective of VI

- Minimize $KL[q(\theta) || p(\theta | D)]$ $-\log -\text{log-lik}$

$$KL[q(\theta) || p(\theta | D)] = -E_{q(\theta)} \left[\log \frac{p(\theta | D)}{q(\theta)} \right]$$

Let's Derive the Objective of VI

- Minimize $KL[q(\theta) || p(\theta | D)]$

$$KL[q(\theta) || p(\theta | D)] = -E_{q(\theta)} \left[\log \frac{p(\theta | D)}{q(\theta)} \right]$$

$$= -E_{q(\theta)} \left[\log \frac{p(\theta, D)}{p(D)q(\theta)} \right] = -E_{q(\theta)} \left[\underbrace{\log \frac{p(\theta, D)}{q(\theta)}} - \underbrace{\log p(D)} \right]$$

Let's Derive the Objective of VI

- Minimize $KL[q(\theta) || p(\theta | D)]$.

$$KL[q(\theta) || p(\theta | D)] = -E_{q(\theta)} \left[\log \frac{p(\theta | D)}{q(\theta)} \right]$$

$$\begin{aligned} &= -E_{q(\theta)} \left[\log \frac{p(\theta, D)}{p(D)q(\theta)} \right] = -E_{q(\theta)} \left[\log \frac{p(\theta, D)}{q(\theta)} - \log p(D) \right] \\ &= \boxed{\log p(D)} - E_{q(\theta)} \left[\log \frac{p(\theta, D)}{q(\theta)} \right] \end{aligned}$$

Model Evidence

sigh

Let's Derive the Objective of VI

Minimize $KL[q(\theta) || p(\theta | D)]$

$$KL[q(\theta) || p(\theta | D)] = \log p(D) - E_{q(\theta)} \left[\log \frac{p(\theta, D)}{q(\theta)} \right]$$

Maximize $E_{q(\theta)} \left[\log \frac{p(\theta, D)}{q(\theta)} \right]$

Let's Derive the Objective of VI

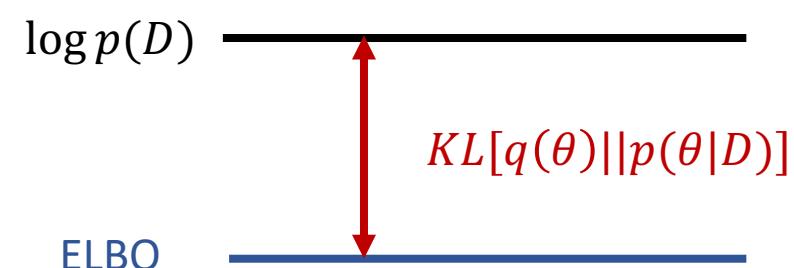
Minimize $KL[q(\theta)||p(\theta|D)]$

$$KL[q(\theta)||p(\theta|D)] = \log p(D) - E_{q(\theta)} \left[\log \frac{p(\theta, D)}{q(\theta)} \right]$$

$$E + \underbrace{KL}_{\gamma^0} = \log p(D) \quad \text{Model Evidence} \Rightarrow \underbrace{\log p(D)}_{} > \underline{E}.$$

Maximize $L = E_{q(\theta)} \left[\log \frac{p(\theta, D)}{q(\theta)} \right]$

Evidence Lower Bound (ELBO)



"Model Evidence = ELBO + KL"

Variational Inference (VI)

$$KL : \underline{q(\theta)} \parallel \underline{p(\theta|D)}$$

The posterior

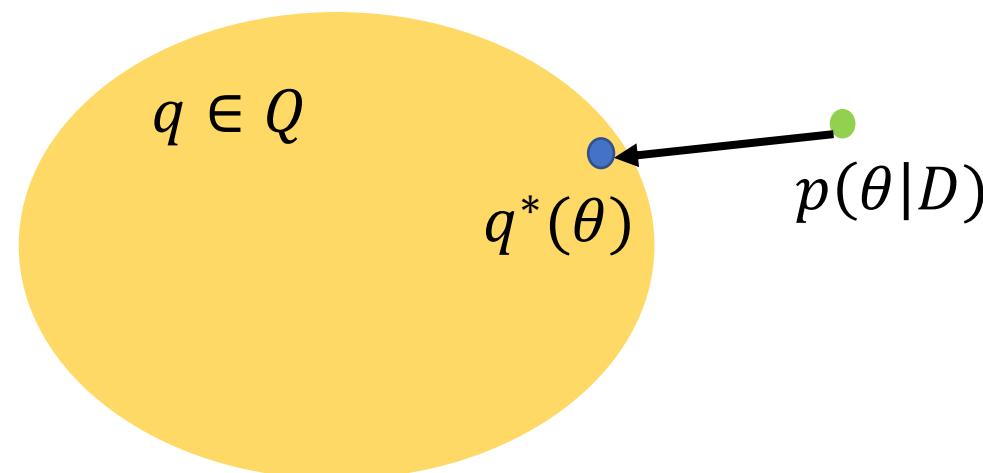
$$p(\theta|D) = p(D|\theta)p(\theta)/p(D)$$

$$p(D|\theta) = \underline{p(D|\theta)} \cdot \underline{p(\theta)}$$

$$q_\phi(\theta)$$

The variational distribution

$$L = E_{\underline{q_\phi(\theta)}} \left[\log \frac{\underline{p(D|\theta)}}{\underline{q_\phi(\theta)}} \right] = \underline{\log p(D)} - \underline{KL[q_\phi(\theta)||p(\theta)]} ,$$



Variational Inference (VI)

- Rewriting the ELBO:

$\log p(D) \geq L = \underbrace{\max_{\text{evidence}}}_{\text{Data fitting term}} E_{q_\phi(\theta)} [\log p(D|\theta)] - \beta \underbrace{KL[q_\phi(\theta) || p(\theta)]}_{\text{KL regulariser}}$

P. q_ϕ . 算法
 M. G. Stan
 $q_\phi(\theta) \sim p(\theta | D)$
 $q_\phi \rightarrow p(\theta)$
 KL ↓
 不變
 不 overfitting

(Negative) Data fitting term:

- Like the usual DL loss you'll use for training neural networks
- ...except that now the network's weights are sampled from q

$$p(\theta | D) = \frac{L}{T} \times \text{prior}$$

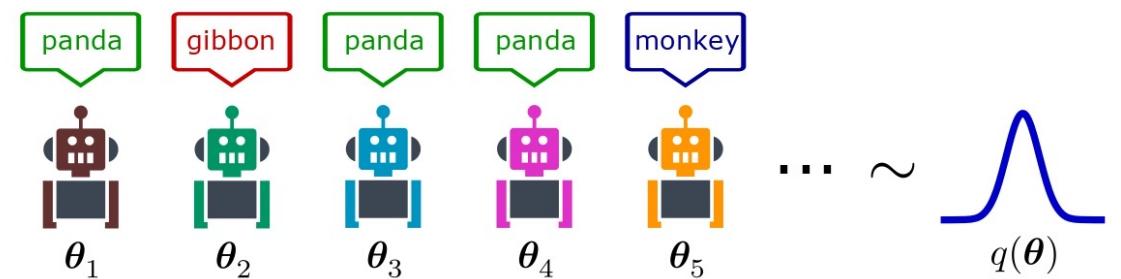
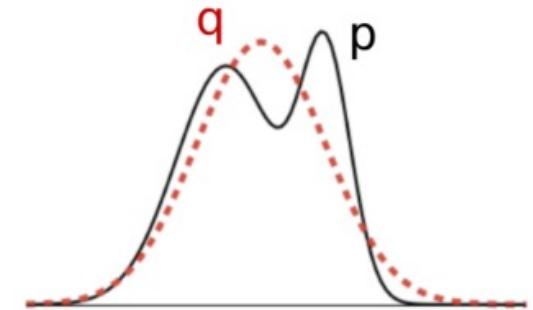
\uparrow
 MLE.

KL regulariser:

- Make the q distribution closer to the prior
- Regularises the approximate posterior, especially when using e.g., Gaussian prior

Approximate Inference in BNNs

- Key steps of approximate inference in BNNs
 1. Construct the $q(\theta) \approx p(\theta | D)$ distribution
 - Simple distributions: e.g. Mean-field Gaussian
 - Structured approximations, e.g. low-rank Gaussians
 - Others (non-Gaussian)
 2. Fit the $q(\theta)$ distribution
 - E.g. with variational inference
 3. Compute prediction with Monte Carlo approximations

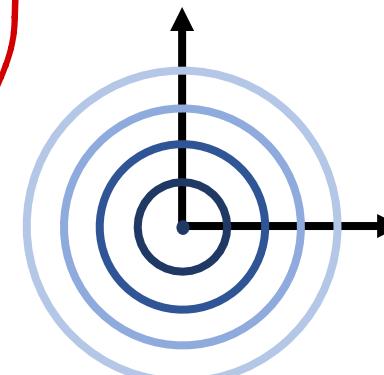


Approximate Inference in BNNs

- Step 1: construct the $q(\theta) \approx p(\theta | D)$ distribution
 - Example: Mean-field Gaussian distribution:

$$q(\theta) = \prod_{l=1}^L q(W^l) q(b^l)$$
$$q(W_l) = \prod_{ij} q(W_{ij}^l), \quad q(W_{ij}^l) = N(W_{ij}^l; M_{ij}^l, V_{ij}^l)$$
$$q(b^l) = \prod_i q(b_i^l), \quad q(b_i^l) = N(b_i^l; m_i^l, v_i^l)$$

STD.



$\rightarrow \mathcal{Z}$

$R . \text{SdR}$

- Variational parameters: $\phi = \{M_{ij}^l, \log V_{ij}^l, m_i^l, \log v_i^l\}_{l=1}^L$

Approximate Inference in BNNs

- Step 2: fit the $q(\theta)$ distribution:

- Variational inference: $\phi^* = \operatorname{argmax} L(\phi)$

$$\underline{L(\phi)} = \underline{E_{q_\phi(\theta)}[\log p(D | \theta)] - KL[q_\phi(\theta) \| p(\theta)]}$$

Approximate Inference in BNNs

- Step 2: fit the $q(\theta)$ distribution:

- Variational inference: $\phi^* = \operatorname{argmax} L(\phi)$ *Likelihood*

$$L(\phi) = \underbrace{E_{q_\phi(\theta)}[\log p(D | \theta)]}_{\text{Likelihood}} - KL[q_\phi(\theta) \| p(\theta)]$$

- First scalable technique: **Stochastic optimization** ←

- i.i.d. assumption of data: $\log p(D | \theta) = \sum_{n=1}^N \log p(y_n | x_n, \theta)$

- Enable mini-batch training with $\{(x_m, y_m)\} \sim D^M$:

$$L(\phi) \approx \frac{1}{M} \sum_{m=1}^M E_{q(\theta)}[\log p(y_m | x_m, \theta)] - KL[q(\theta) \| p(\theta)]$$

$$w, b \sim \mathcal{N}_{V_w}^M, b \sim \mathcal{N}_{V_b}^M$$

←

Approximate Inference in BNNs

- Step 2: fit the $q(\theta)$ distribution:

- Variational inference: $\phi^* = \operatorname{argmax} L(\phi)$

$$L(\phi) = E_{q_\phi(\theta)}[\log p(D | \theta)] - KL[q_\phi(\theta) \| p(\theta)]$$

- First scalable technique: **Stochastic optimization**

- i.i.d. assumption of data: $\log p(D | \theta) = \sum_{n=1}^N \log p(y_n | x_n, \theta)$

- Enable mini-batch training with $\{(x_m, y_m)\} \sim D^M$:

$$L(\phi) \approx \frac{N}{M} \sum_{m=1}^M E_{q(\theta)}[\log p(y_m | x_m, \theta)] - KL[q(\theta) \| p(\theta)]$$

rewighting to ensure calibrated
posterior concentration

Approximate Inference in BNNs

- Step 2: fit the $q(\theta)$ distribution:
 - 2nd scalable technique: **Monte Carlo sampling**
 - $E_{q(\theta)}[\log p(y | x, \theta)]$ intractable even with Gaussian $q(\theta)$
 - Solution: Monte Carlo estimate:

$$E_{q(\theta)}[\log p(y | x, \theta)] \approx \frac{1}{K} \sum_k^K \log p(y | x, \theta_k), \quad \theta_k \sim q(\theta)$$

Approximate Inference in BNNs

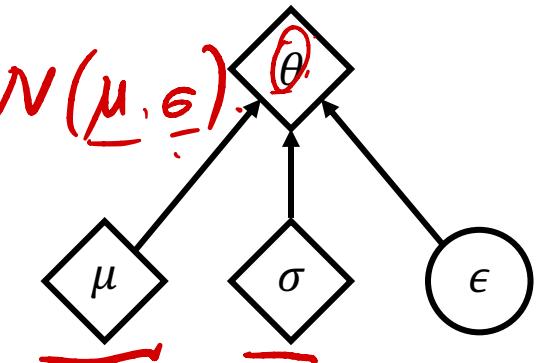
- Step 2: fit the $q(\theta)$ distribution:
 - 2nd scalable technique: **Monte Carlo sampling**
 - $E_{q(\theta)}[\log p(y | x, \theta)]$ intractable even with Gaussian $q(\theta)$
 - Solution: Monte Carlo estimate:

$$E_{q(\theta)}[\log p(y | x, \theta)] \approx \frac{1}{K} \sum_k^K \log p(y | x, \theta_k), \quad \theta_k \sim q(\theta)$$

• Reparameterization trick to sample mean-field Gaussians:

$$\underbrace{\theta_k \sim q(\theta)}_{\text{DK.}} \Leftrightarrow \underbrace{\theta_k = m_\theta + \sigma_\theta \odot \epsilon_k}_{\text{!}} \quad \underbrace{\epsilon_k \sim N(0, I)}$$

$$\dots \quad \frac{\partial \theta_k}{\partial m_\theta} \quad \frac{\partial \theta_k}{\partial \sigma_\theta}$$



Approximate Inference in BNNs

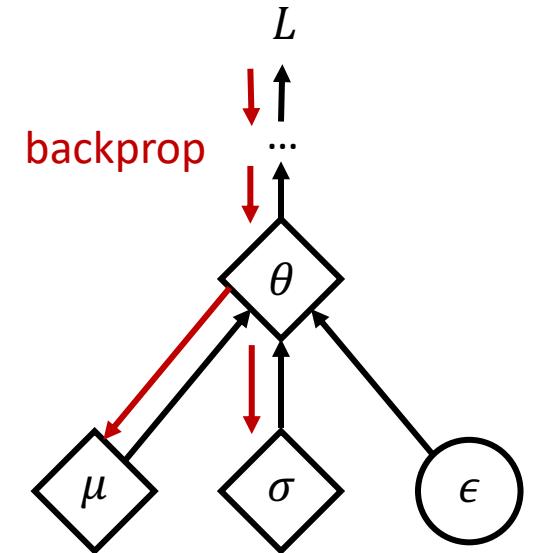
- Step 2: fit the $q(\theta)$ distribution:
 - 2nd scalable technique: **Monte Carlo sampling**
 - $E_{q(\theta)}[\log p(y | x, \theta)]$ intractable even with Gaussian $q(\theta)$
 - **Solution: Monte Carlo estimate:**

$$E_{q(\theta)}[\log p(y | x, \theta)] \approx \frac{1}{K} \sum_k^K \log p(y | x, \theta_k), \quad \theta_k \sim q(\theta)$$

- **Reparameterization trick** to sample mean-field Gaussians:

$$\theta_k \sim q(\theta) \Leftrightarrow \theta_k = m_\theta + \sigma_\theta \odot \epsilon_k, \quad \epsilon_k \sim N(0, I)$$

$$\Rightarrow E_{q(\theta)} [\log p(y | x, \theta)] \approx \frac{1}{K} \sum_k^K \log p(y | x, \theta_k = m_\theta + \sigma_\theta \epsilon_k), \quad \epsilon_k \sim N(0, I)$$



Approximate Inference in BNNs

- Combining both steps:

$$L(\phi) \approx \frac{N}{M} \sum_{m=1}^M \frac{1}{K} \sum_{k=1}^K \log p(y_m | x_m, \theta_k) - \underline{KL[q(\theta) || p(\theta)]}, \theta_k \sim q(\theta)$$

ϕ p.

(if not, can also be estimated with Monte Carlo)

analytic between two Gaussians

In regression:

$$p(y | x, \theta) = N(f_\theta(x), \sigma^2)$$

In classification:

$$p(y | x, \theta) = \text{Categorical}(\text{logit} = f_\theta(x))$$

Approximate Inference in BNNs

MC.

- Step 3: compute prediction with Monte Carlo approximations:

$$\underbrace{p(y^* | x^*, D)}_{\text{Step 3}} \approx \frac{1}{K} \sum_{k=1}^K p(y^* | x^*, \underline{\theta_k}), \quad \underline{\theta_k \sim q(\theta)} \quad \underbrace{p(\theta | D)}_{\text{MC.}}$$

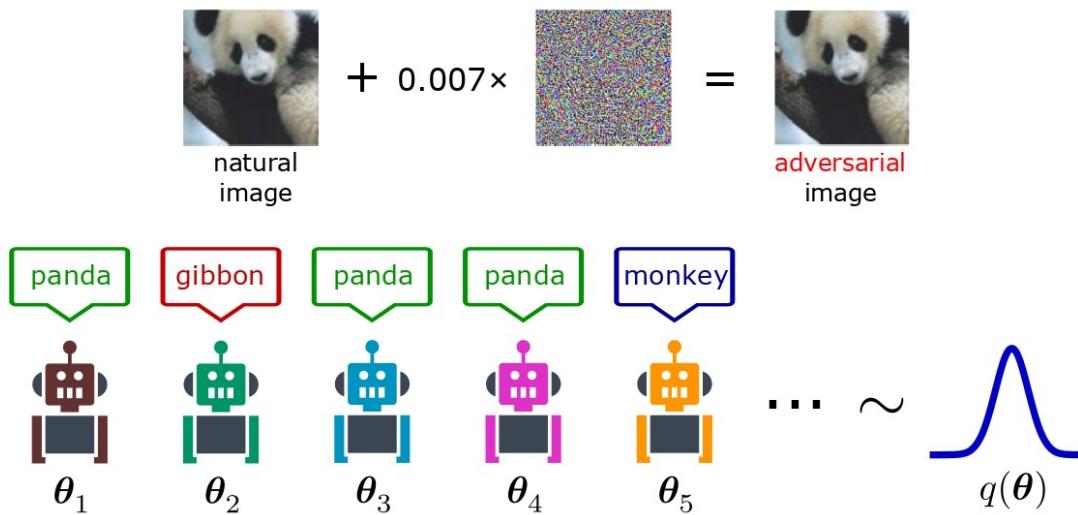
Mean-field Gaussian case:

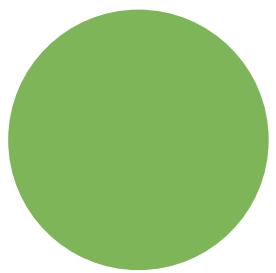
$$\theta_k = \underbrace{m_\theta + \sigma_\theta \odot \epsilon_k}_{\text{mean } \theta}, \quad \epsilon_k \sim N(0, I)$$

$$\sigma_\theta$$

$$g(\theta) \\ g(w) \sim \mathcal{N}(M, V).$$

$$w = M + V \cdot \underline{\epsilon} \\ \epsilon \sim N(0, 1)$$





Part II: Bayesian MLPs

- **Implement various BNN methods for MLP architectures**
- **Regression example test**
- **Case study 1: Bayesian Optimisation with UCB**

<https://bit.ly/3zF1zvA>

Instructions for using this Google Colab notebook:

- Make sure you have signed in with your Google account;
- Click “File > Save a copy in Drive” to create your own copy;
- Let’s play around with the demo using your own copy!

Variational Infer - Findings with MFVI for Bayesian MLPs

mean field $\Rightarrow q(\theta) = \underbrace{q(w)}_{\prod q(w_i)}$ $\cdot q(b)$.

Findings with MFVI for Bayesian MLPs

- MFVI tends to underfit
 - Initialisation matters ←
 - Tuning the beta parameter also helps

$$W_b \leftarrow \underline{m_b, v_b} \quad \text{prior.}$$



$$\text{ELBO} = \text{data fitting} - \beta \text{KL}[q_{\phi}(\theta) || p(\theta)]$$

No free lunch theorem.

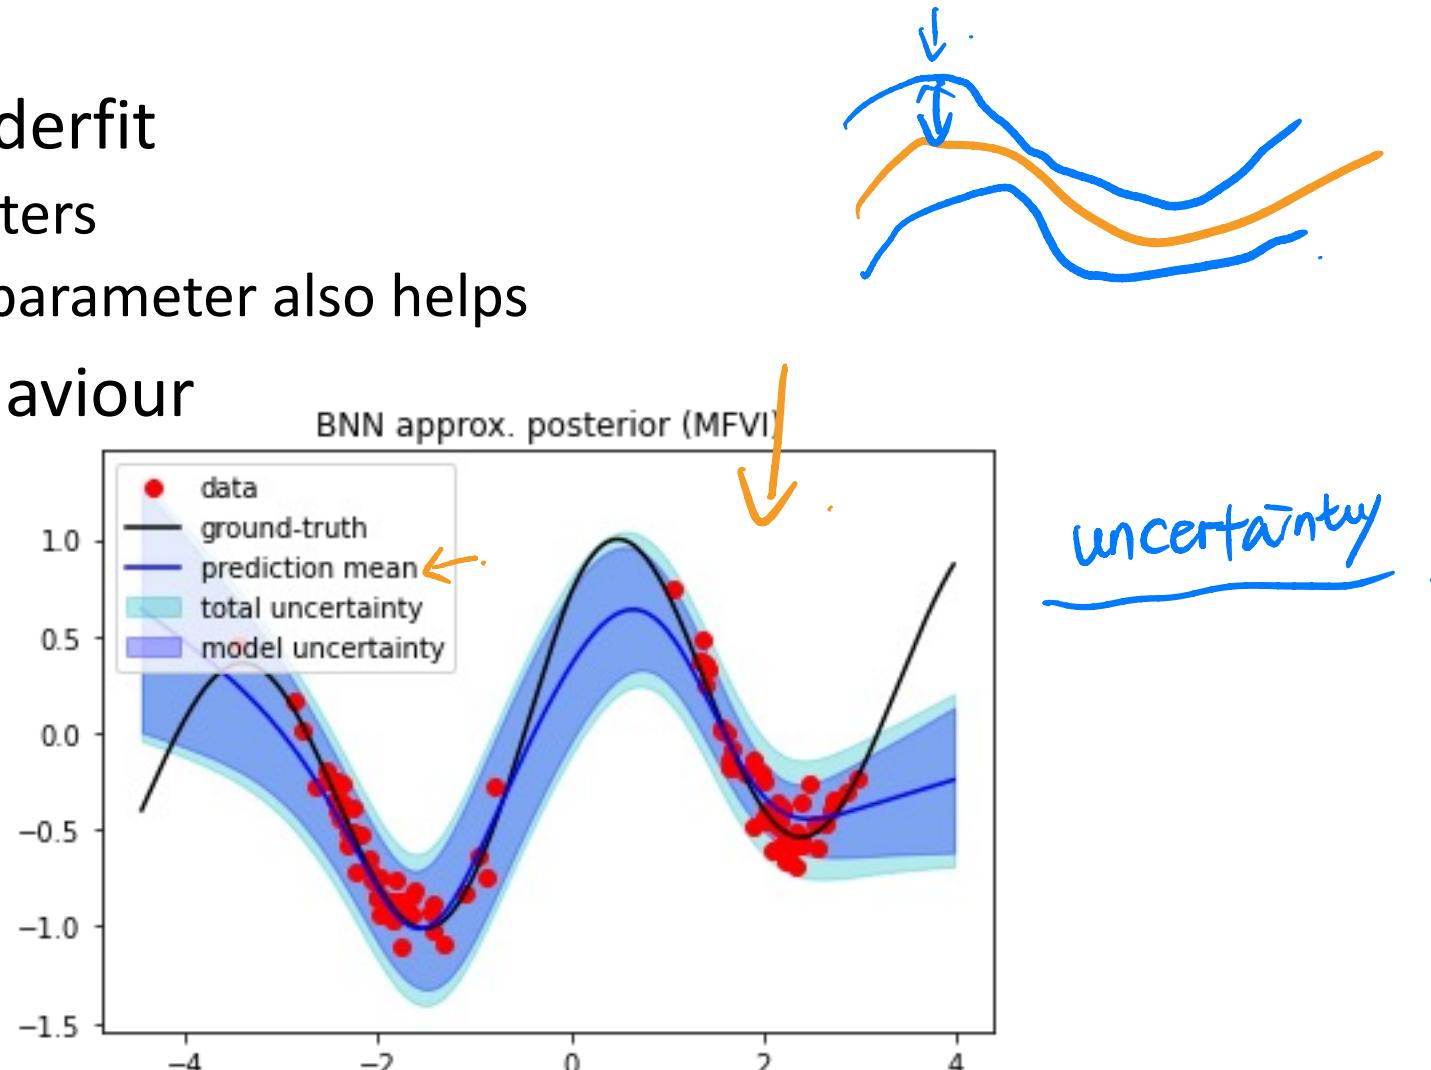
没有一个模型可以在所有任务上都表现得很好。

模型 ← 内含先验。

bayes: 显式的先验 $p(\theta)$. 没有一个通用的最好的 prior.

Findings with MFVI for Bayesian MLPs

- MFVI tends to underfit
 - Initialisation matters
 - Tuning the beta parameter also helps
- Uncertainty behaviour



Using other q distributions?

- Using more complicated q distributions? q : mean field.
 - Pros: more flexible approximations \Rightarrow better posterior approximations (?)
 - Cons: higher time & space complexities

$$q(\underline{w}) = \prod q(w_{ij}^t),$$
$$\underline{q}_b(\underline{\underline{w}}) = \underline{\underline{[}} \underline{\underline{]}).$$

full

Using other q distributions?

- Using more complicated q distributions?
 - Pros: more flexible approximations \Rightarrow better posterior approximations (?)
 - Cons: higher time & space complexities
 - We will look at 2 alternatives:
 - “Last-layer BNN”: Full covariance Gaussian approximations for the last layer
 - MC-Dropout: adding dropout layers and run them in both train & test time
- deterministic*
- $0 \rightarrow$ random variabb.

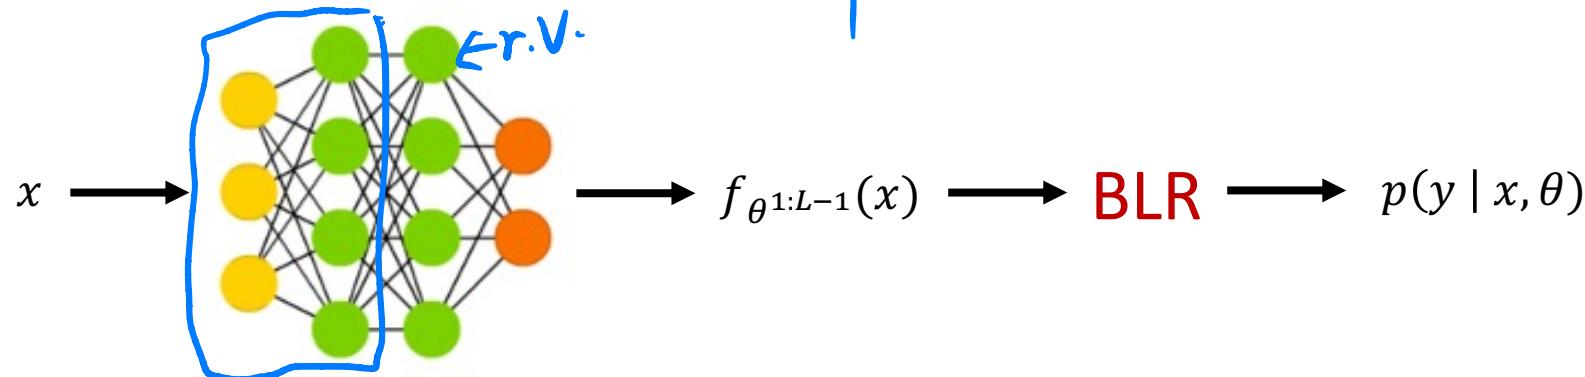
“Last-layer BNN”

Layer.

- Use deterministic layers for all but the last layer
- For the last layer: Use Full-covariance Gaussian approximate posterior:

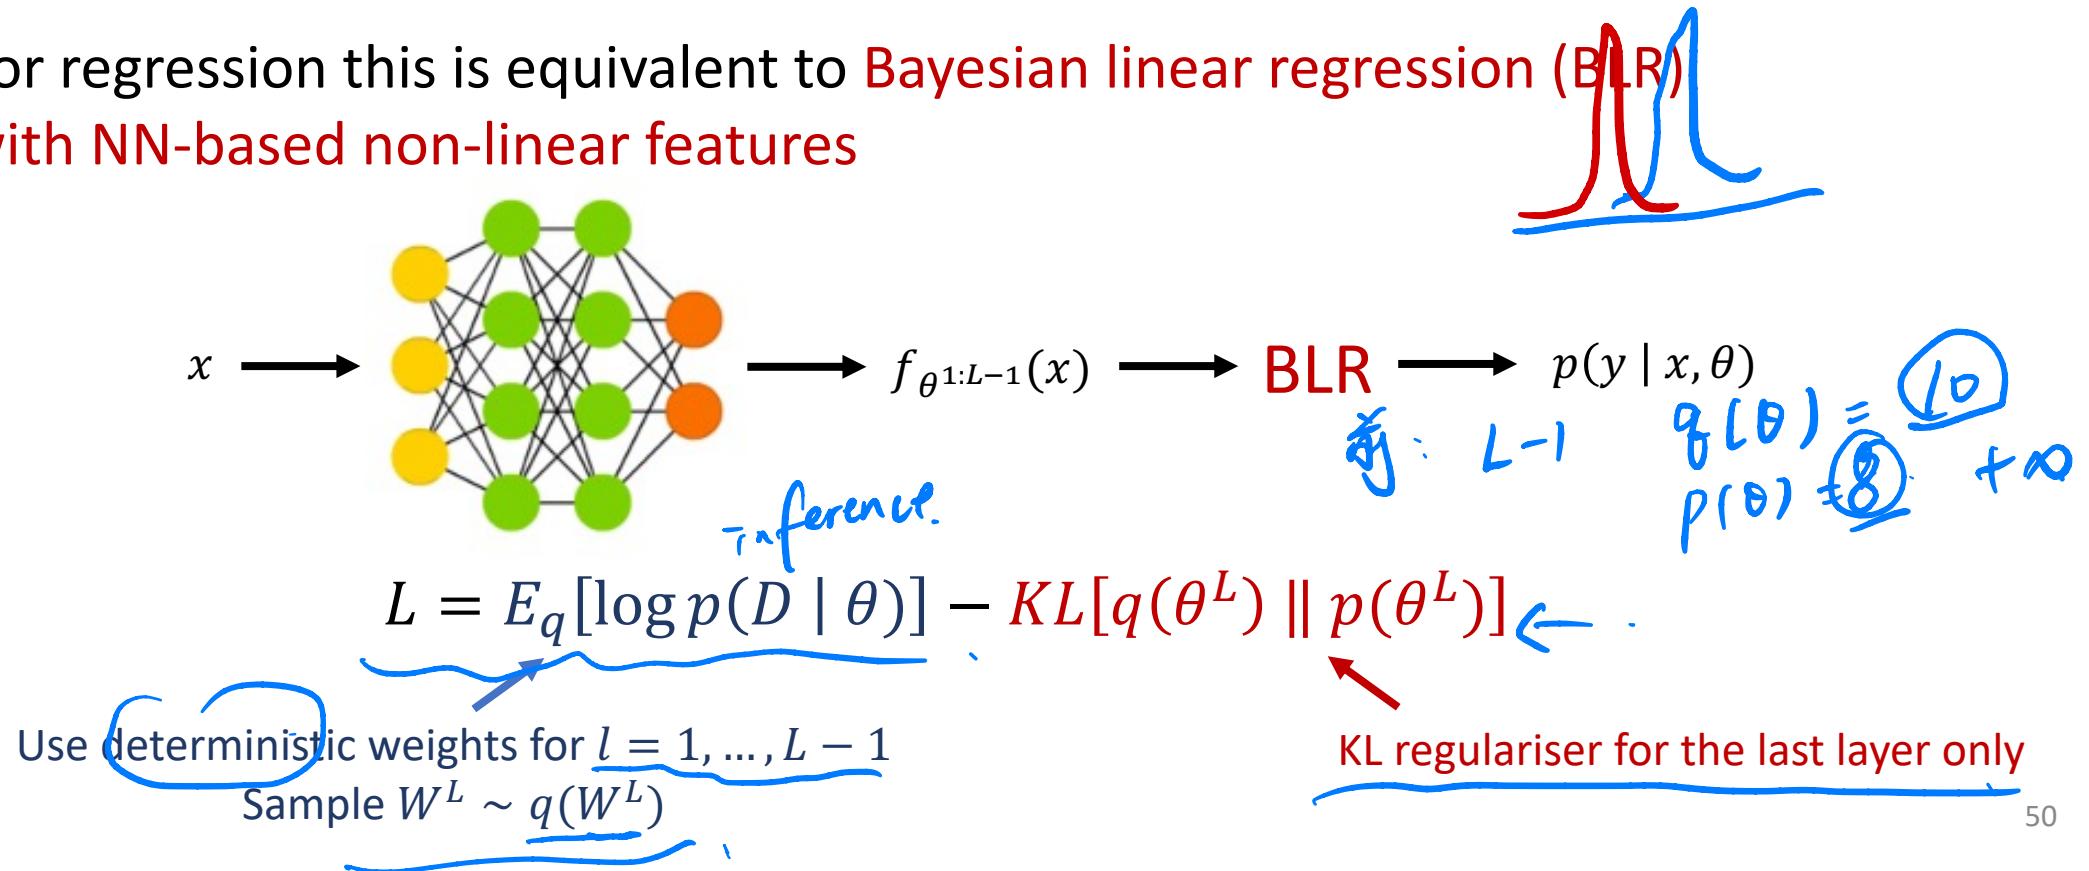
$$\left\{ \begin{array}{l} q(\theta^l) = \delta(W^l = M^l, b^l = m^l), l = 1, \dots, L-1, \\ q(\theta^L) = N(vec(\theta^L); vec(\mu^L), \Sigma), \theta^L = \{W^L, b^L\} \end{array} \right.$$

- For regression this is equivalent to **Bayesian linear regression (BLR)** with NN-based non-linear features



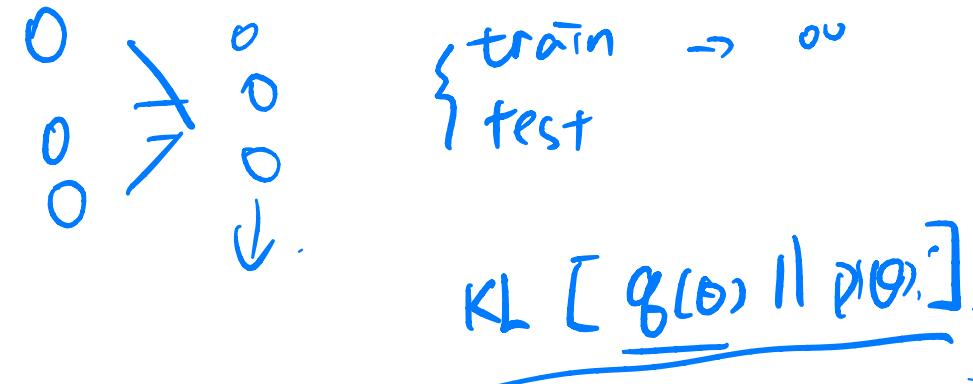
“Last-layer BNN”

- Use deterministic layers for all but the last layer
- For the last layer: Use Full-covariance Gaussian approximate posterior
- For regression this is equivalent to **Bayesian linear regression (BLR)** with NN-based non-linear features



MC-Dropout

- Add dropout layers to the network
- Perform dropout during training



$$L = \mathbb{E}_q [\log p(D|\theta)] - (1 - \pi) \ell_2(\phi)$$

The MC sampling procedure is implicitly defined

L2 regulariser on the variational parameters

KL.

$W_0 \sim \underline{\mathcal{B}(0, \sigma)}$

$w_0 \sim N(0/\sigma)$

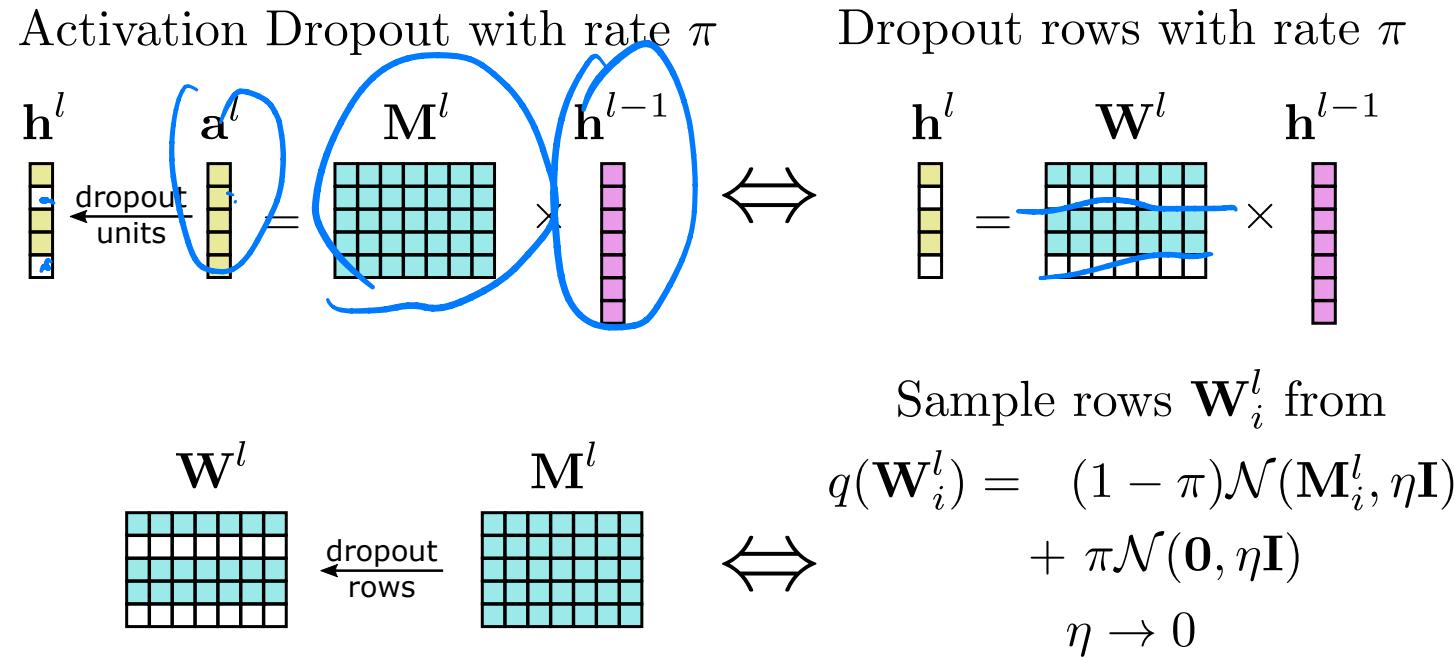
- In test time, run multiple forward passes with dropout

$$p(y^* | x^*, D) \approx \frac{1}{K} \sum_{k=1}^K p(y^* | x^*, \theta_k), \quad \theta_k \sim q(\theta)$$

The MC sampling procedure is implicitly defined

MC-Dropout

- Two equivalent ways to implement MC-Dropout:



(Similar logic applies when including the bias terms, see lecture notes.)

(Notice that pytorch's nn.Linear layer uses formats like xW^T instead of Wx .)

Using other q distributions?

- What you'll do for the next part of the tutorial:
 - Implement MC-Dropout in 2 ways
 - Run the regression sample with the 2 approximation methods discussed
 - Compare with MFVI

Case study 1: Bayesian Optimisation

- Imagine you'd like to solve the following task:

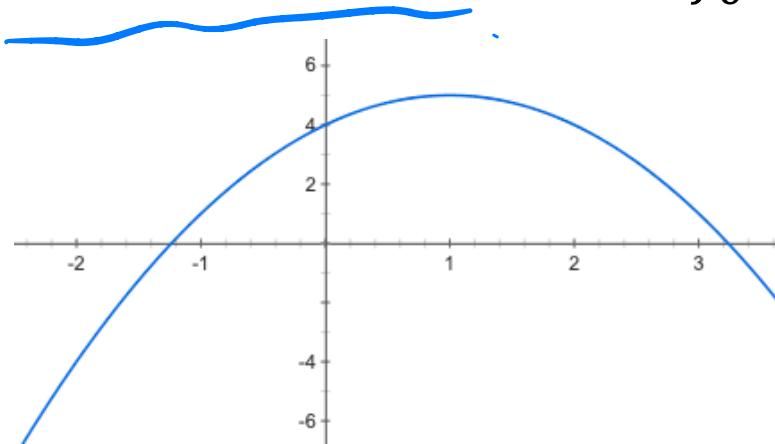
$$\underline{x^*} = \operatorname{argmax}_x f_0(x) \text{ 等价.}$$

Case study 1: Bayesian Optimisation

- Imagine you'd like to solve the following task:

$$x^* = \operatorname{argmax}_x f_0(x) \quad \text{gradient}$$

Known functional form of f_0 :



Gradient descent, Newton's method,

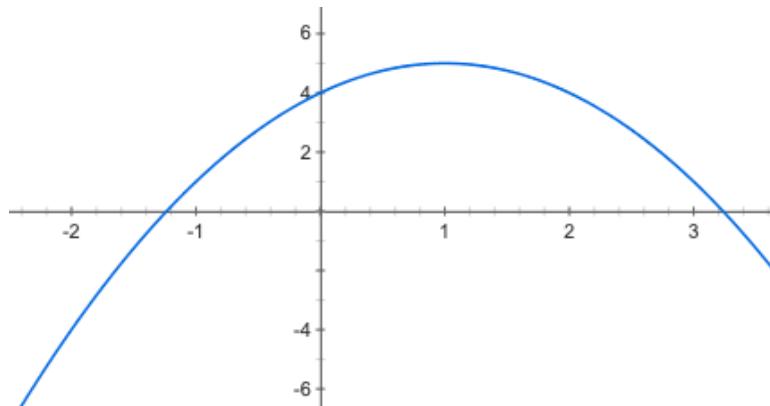
...

Case study 1: Bayesian Optimisation

- Imagine you'd like to solve the following task:

$$x^* = \operatorname{argmax}_x f_0(x)$$

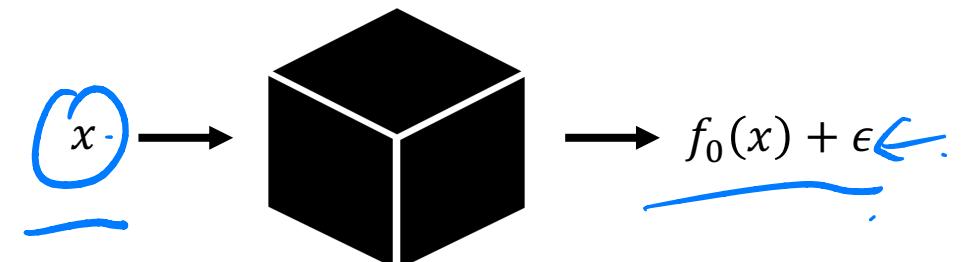
Known functional form of f_0 :



Gradient descent, Newton's method,

...

Unknown functional form of f_0 :

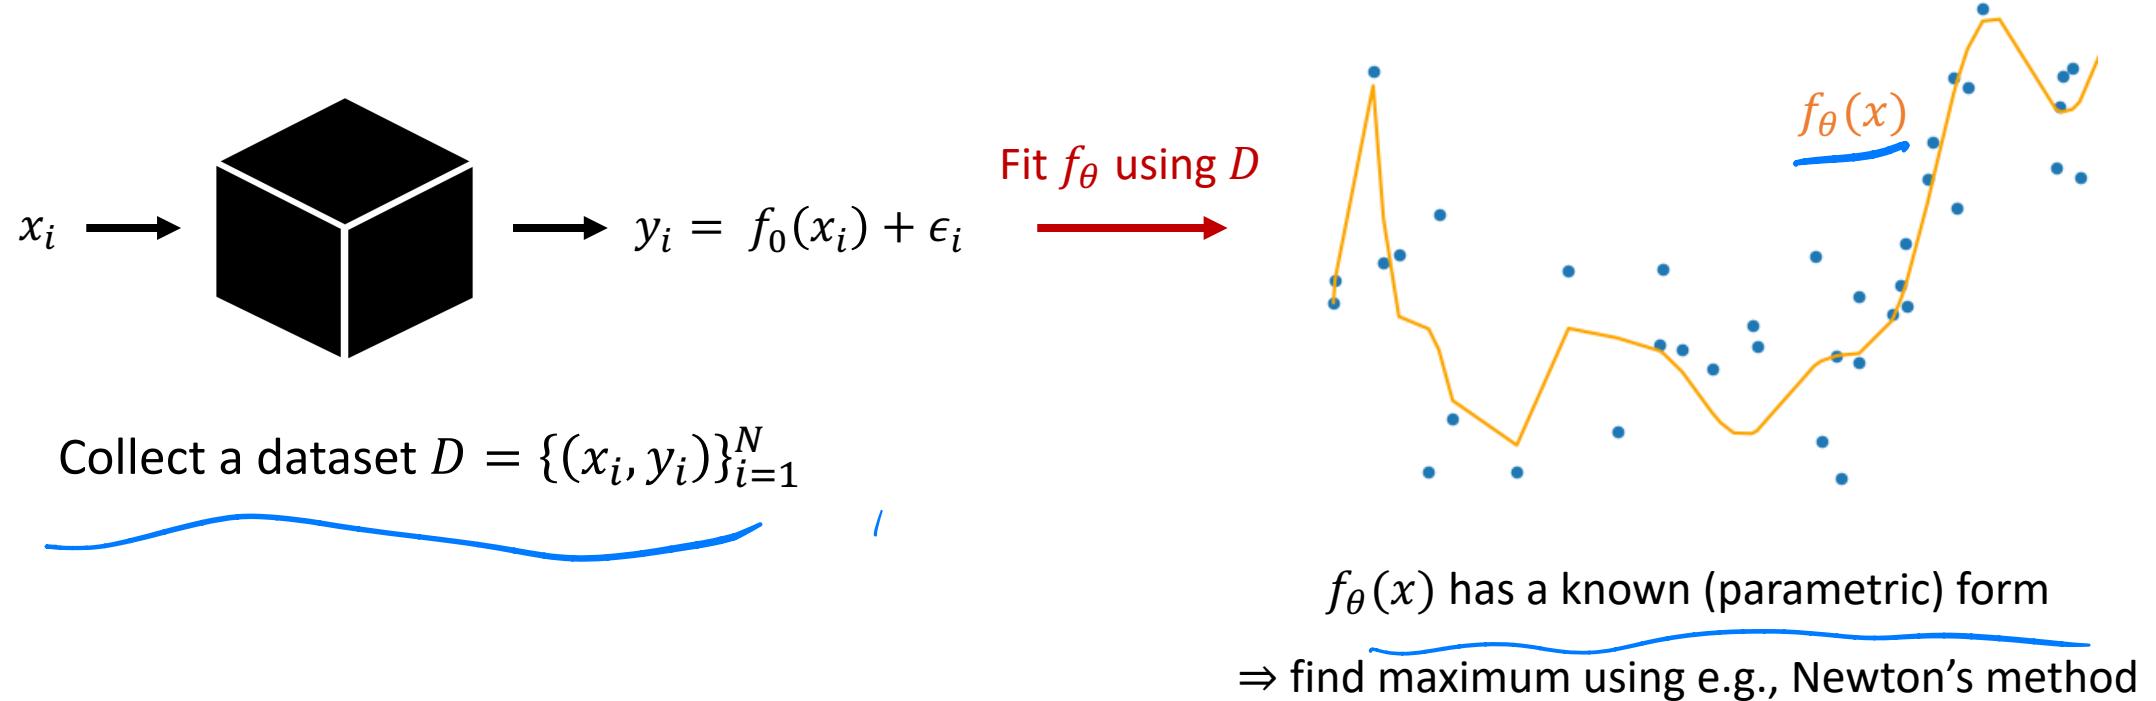


(can only query (noisy) function values)



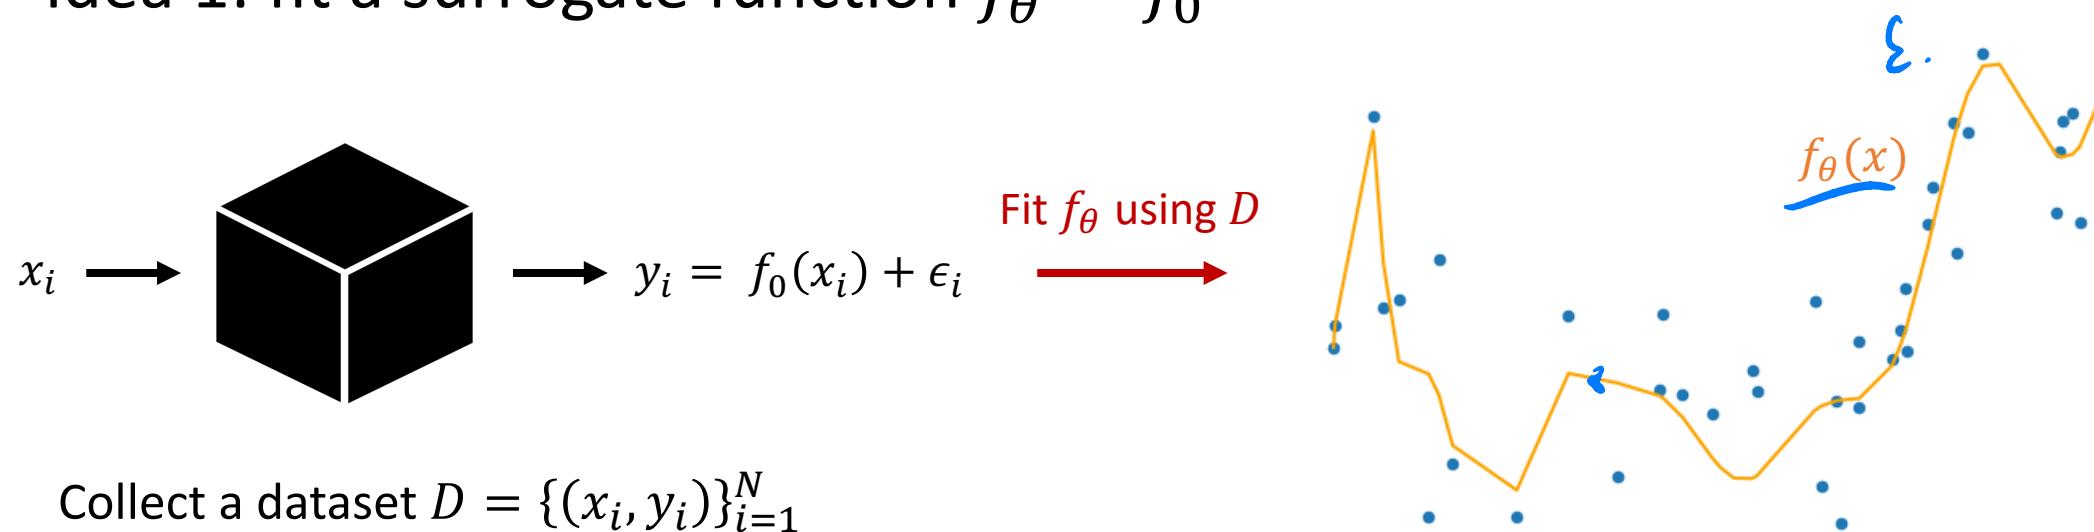
Case study 1: Bayesian Optimisation

- Idea 1: fit a surrogate function $f_\theta \approx f_0$



Case study 1: Bayesian Optimisation

- Idea 1: fit a surrogate function $f_\theta \approx f_0$

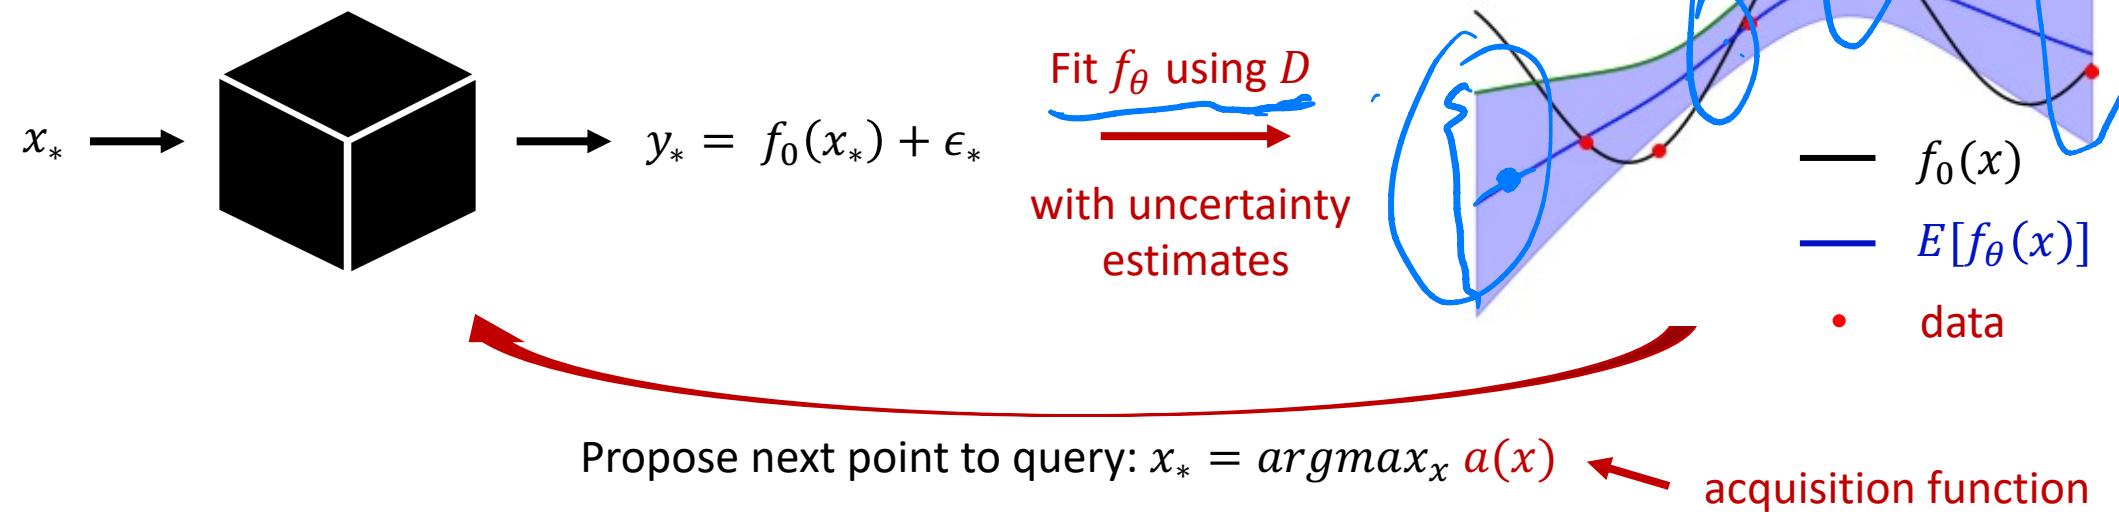


- Issues of this approach:
 - Need to collect a lot of datapoints for accurate fitting of f_θ
 - Do not consider uncertainty at unseen locations

Case study 1: Bayesian Optimisation

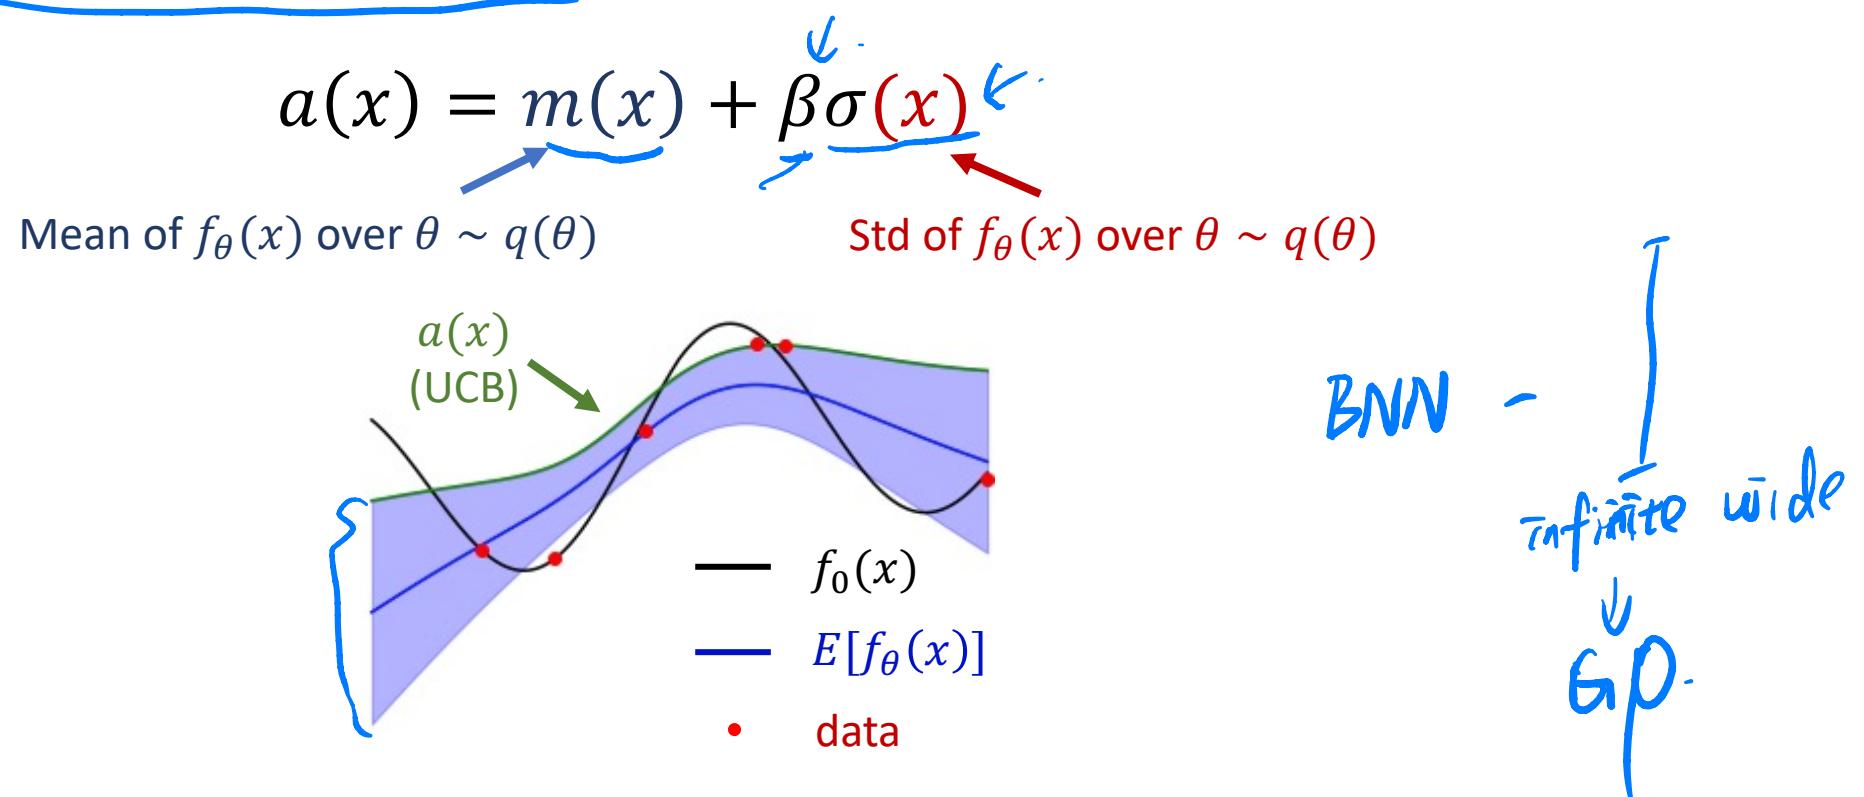
- Idea of BO: iterate the following steps
 - fit a surrogate function f_θ with uncertainty estimates
 - Use the surrogate function to guide the dataset collection process

Update the dataset $D = D \cup \{(x_*, y_*)\}$



Case study 1: Bayesian Optimisation

- Upper confidence bound (UCB): a widely used acquisition function

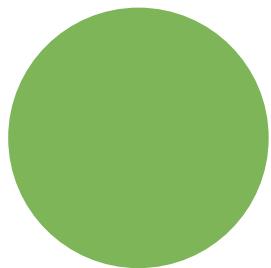


Case study 1: Bayesian Optimisation

- What you'll do for the case study part of the tutorial:
 - Implement UCB acquisition function
 - Run the BO example
 - Play around with hyper-parameters and other settings

Example answers of the tutorial demos:

Regression: <https://bit.ly/39eZHit>



Part III: Bayesian ConvNets

- Classification example test
- Case study 2: Detecting adversarial examples

<https://bit.ly/3Hd1Ass>

Instructions for using this Google Colab notebook:

- Make sure you have signed in with your Google account;
- Click “File > Save a copy in Drive” to create your own copy;
- Use GPU: in “Runtime > Change runtime type”, choose “GPU” for “hardware accelerator”
- Let’s play around with the demo using your own copy!

Case study 2: Detecting adversarial examples

□ → BNN → (lable, 置信度). ↓

- Hypothesis:
 - Adversarial examples are regarded as OOD data 非 OOD, noisy, 假数据
 - BNNs become uncertain about their prediction on OOD data
 - ⇒ uncertainty measures can be used for detecting adversarial examples

$$\frac{y = y^* + \epsilon}{\hat{y} = x^* + \epsilon}$$

(y, x)

Uncertainty measures

Total uncertainty = epistemic uncertainty + aleatoric uncertainty

Due to lack of "knowledge"
(reducible when having more data)



数据

固有
不确定性

Due to inherent stochasticity in data
(non-reducible)

Imagine flipping a coin:

- Epistemic uncertainty: "How much do I believe the coin is fair?"
 - Model's belief after seeing the population
 - Reduces when having more data
- Aleatoric uncertainty: "What's the next coin flip outcome?"
 - Individual experiment outcome
 - Non-reducible



Uncertainty measures

Total uncertainty = **epistemic uncertainty** + **aleatoric uncertainty**

Due to lack of “knowledge”
(reducible when having more data)

Due to inherent stochasticity in data
(non-reducible)

Computing uncertainty in classification models:

C 素例.

$$H[p] = - \sum_{c=1}^C p_c \log p_c,$$

$$p = (p_1, \dots, p_C), \sum_{c=1}^C p_c = 1$$

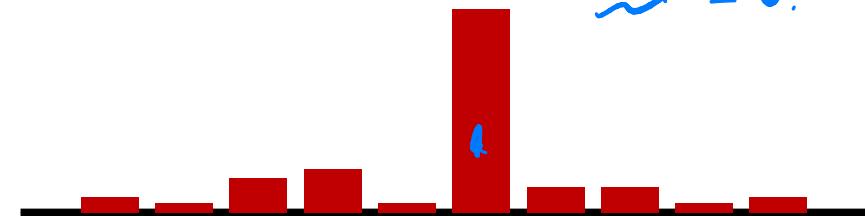
High entropy: $H[p] \rightarrow \log C$

$$p_c = \frac{1}{C}$$



Low entropy: $H[p] \rightarrow 0$

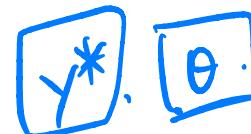
$$p_c = 1 \quad \sim = 0.$$



Uncertainty measures

Total uncertainty = **epistemic uncertainty** + **aleatoric uncertainty**

Due to lack of “knowledge”
(reducible when having more data)



Due to inherent stochasticity in data
(non-reducible)

Computing uncertainty in classification models:

Recall for Bayesian predictive distribution:

(y^*, x^*)

$$p(y^* | x^*, D) = \int p(y^* | x^*, \theta) p(\theta | D) d\theta$$

$$H[y^* | x^*, D] = I[y^* ; \theta | x^*, D] + E_{p(\theta | D)} [H[y^* | x^*, \theta]]$$

Total entropy of the
predictive distribution

Mutual information
between y^* and θ

Conditional entropy
under posterior

數學
 y^*, θ 的互信息。

Uncertainty measures

Total uncertainty = **epistemic uncertainty** + **aleatoric uncertainty**

Due to lack of “knowledge”
(reducible when having more data)

Due to inherent stochasticity in data
(non-reducible)

Computing uncertainty in classification models:

Recall for Bayesian predictive distribution **with approximation**:

$$\underbrace{p(y^* | x^*, D)}_{\text{Total entropy (for total uncertainty)}} \approx \frac{1}{K} \sum_{k=1}^K p(y^* | x^*, \theta_k), \quad \theta_k \sim q(\theta).$$

Total entropy (for total uncertainty):

$$H[y^* | x^*, D] \approx H[\underbrace{\frac{1}{K} \sum_{k=1}^K p(y^* | x^*, \theta_k)}_{\text{Total entropy}}]$$

Uncertainty measures

Total uncertainty = **epistemic uncertainty** + **aleatoric uncertainty**

Due to lack of “knowledge”
(reducible when having more data)

Due to inherent stochasticity in data
(non-reducible)

Computing uncertainty in classification models:

Recall for Bayesian predictive distribution **with approximation**:

$$p(y^* | x^*, D) \approx \frac{1}{K} \sum_{k=1}^K p(y^* | x^*, \theta_k), \quad \theta_k \sim q(\theta)$$

Conditional entropy (for aleatoric uncertainty): Intuition.

$$\underbrace{E_{p(\theta | D)}[H[y^* | x^*, \theta]]}_{\text{Intuition}} \approx \frac{1}{K} \sum_{k=1}^K H[p(y^* | x^*, \theta_k)] \quad \theta_k \sim p(\theta | D)$$

Uncertainty measures

Total uncertainty = **epistemic uncertainty** + **aleatoric uncertainty**

Due to lack of “knowledge”
(reducible when having more data)

Due to inherent stochasticity in data
(non-reducible)

Computing uncertainty in classification models:

Recall for Bayesian predictive distribution **with approximation**:

$$p(y^* | x^*, D) \approx \frac{1}{K} \sum_{k=1}^K p(y^* | x^*, \theta_k), \quad \theta_k \sim q(\theta)$$

Mutual information (for epistemic uncertainty):

$$I[y^*; \theta | x^*, D] \approx H\left[\frac{1}{K} \sum_{k=1}^K p(y^* | x^*, \theta_k)\right] - \underbrace{\frac{1}{K} \sum_{k=1}^K H[p(y^* | x^*, \theta_k)]}_{\text{Total}}$$

↓.
↓.

Uncertainty measures

Total uncertainty = **epistemic uncertainty** + **aleatoric uncertainty**

Due to lack of “knowledge”
(reducible when having more data)

Due to inherent stochasticity in data
(non-reducible)

Computing uncertainty in classification models:

Recall for Bayesian predictive distribution **with approximation**:

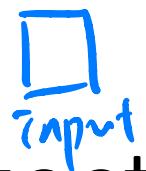
$$p(y^* | x^*, D) \approx \frac{1}{K} \sum_{k=1}^K p(y^* | x^*, \theta_k), \quad \theta_k \sim q(\theta)$$

似然函数 3 次
 x^*, y^*
PLO) 认知改变.

Mutual information (for epistemic uncertainty) if you can do exact inference:

$$\downarrow I[y^*; \theta | x^*, D] = E_{p(y^*|x^*, D)} [KL[p(\theta | D, x^*, y^*) || p(\theta | D)]]$$

↓ “What the model thinks the posterior is going to change if we add new observation at location x^* ”



$\Rightarrow \gamma^*$, uncertainty.

OOD/noisy
/

Case study 2: Detecting adversarial examples

- What you'll do for the case study part of the tutorial:
 - Implement the uncertainty measures
 - Total entropy, conditional entropy, and mutual info
 - Run adversarial attacks on various trained networks
 - See how diversity helps in detecting adversarial examples
 - Detection by thresholding the uncertainty measures
 - We consider best TPR with $FPR \leq 5\%$

Ensemble BNNs

$$\underbrace{p(y^* | x^*)}_{\theta} \cdot p(y^* | x^*, \theta^I).$$

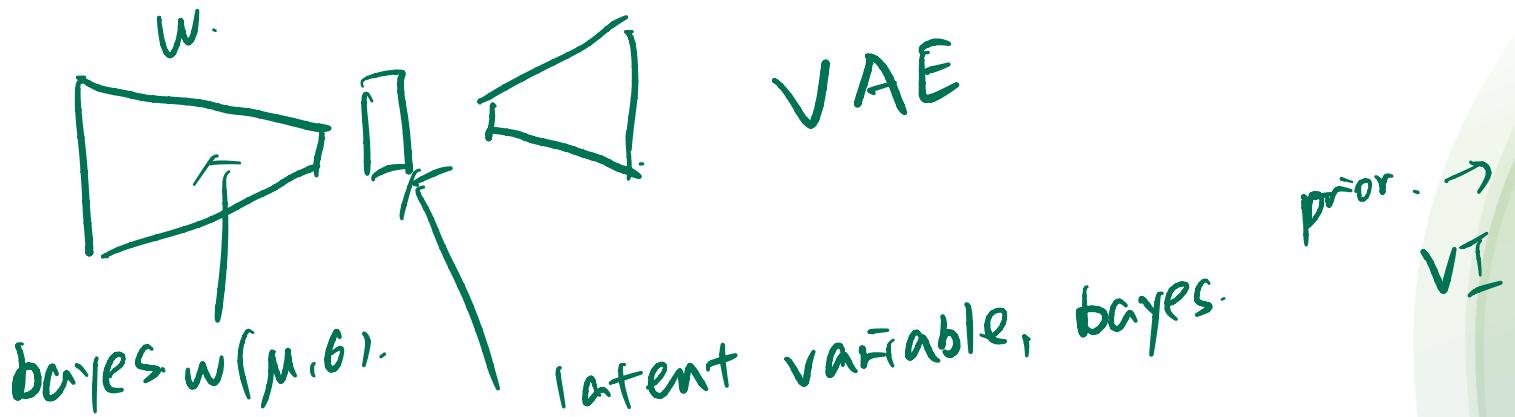
- Define q distribution as mixture of mean-field Gaussian:

$$q(\theta) = \frac{1}{S} \sum_{s=1}^S q(\theta|s), \quad q(\theta|s) = N(\theta; \mu_s, \text{diag}(\sigma_s^2))$$

- Objective is still a valid lower-bound to $\log p(D)$:

$$L = \frac{1}{S} \sum_{s=1}^S ELBO[q(\theta|s)], \quad ELBO[q(\theta|s)] = E_{q(\theta|s)}[\log p(D | \theta)] - KL[q(\theta|s) \| p(\theta)]$$

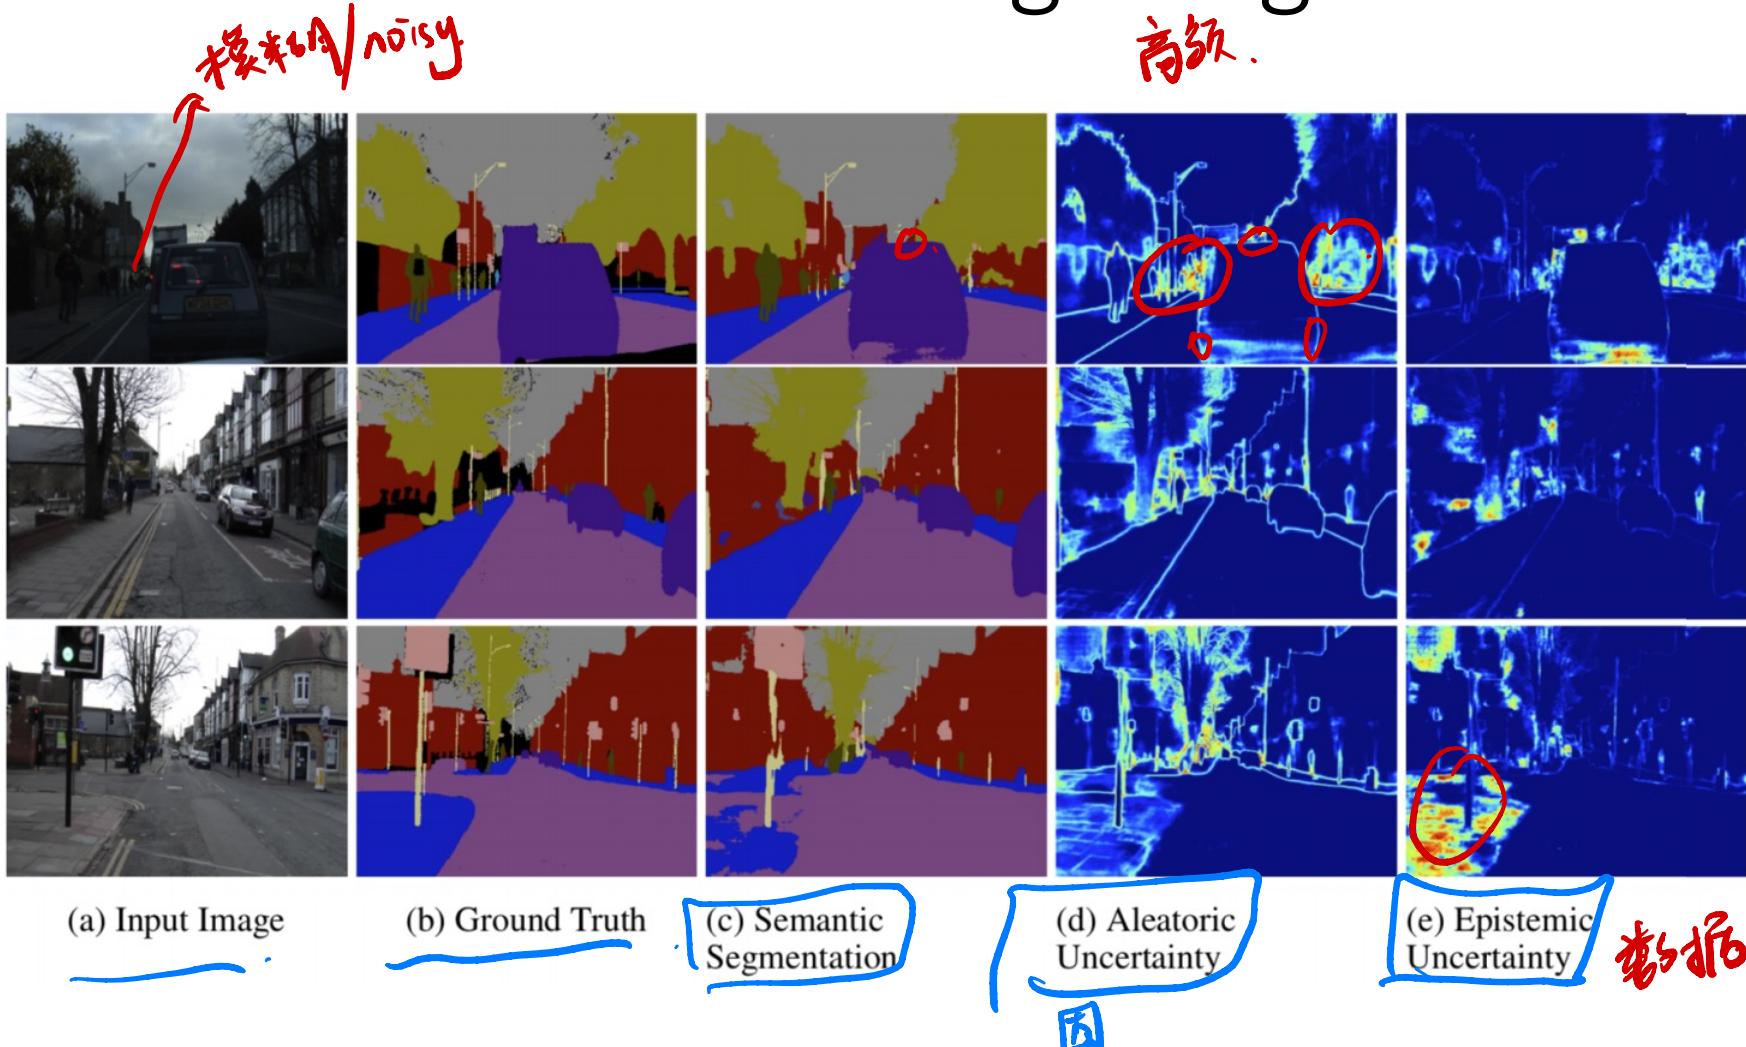
- The parameters of $q(\theta|s)$ for different s are **independent**
⇒ train S number of MFVI-BNNs independently



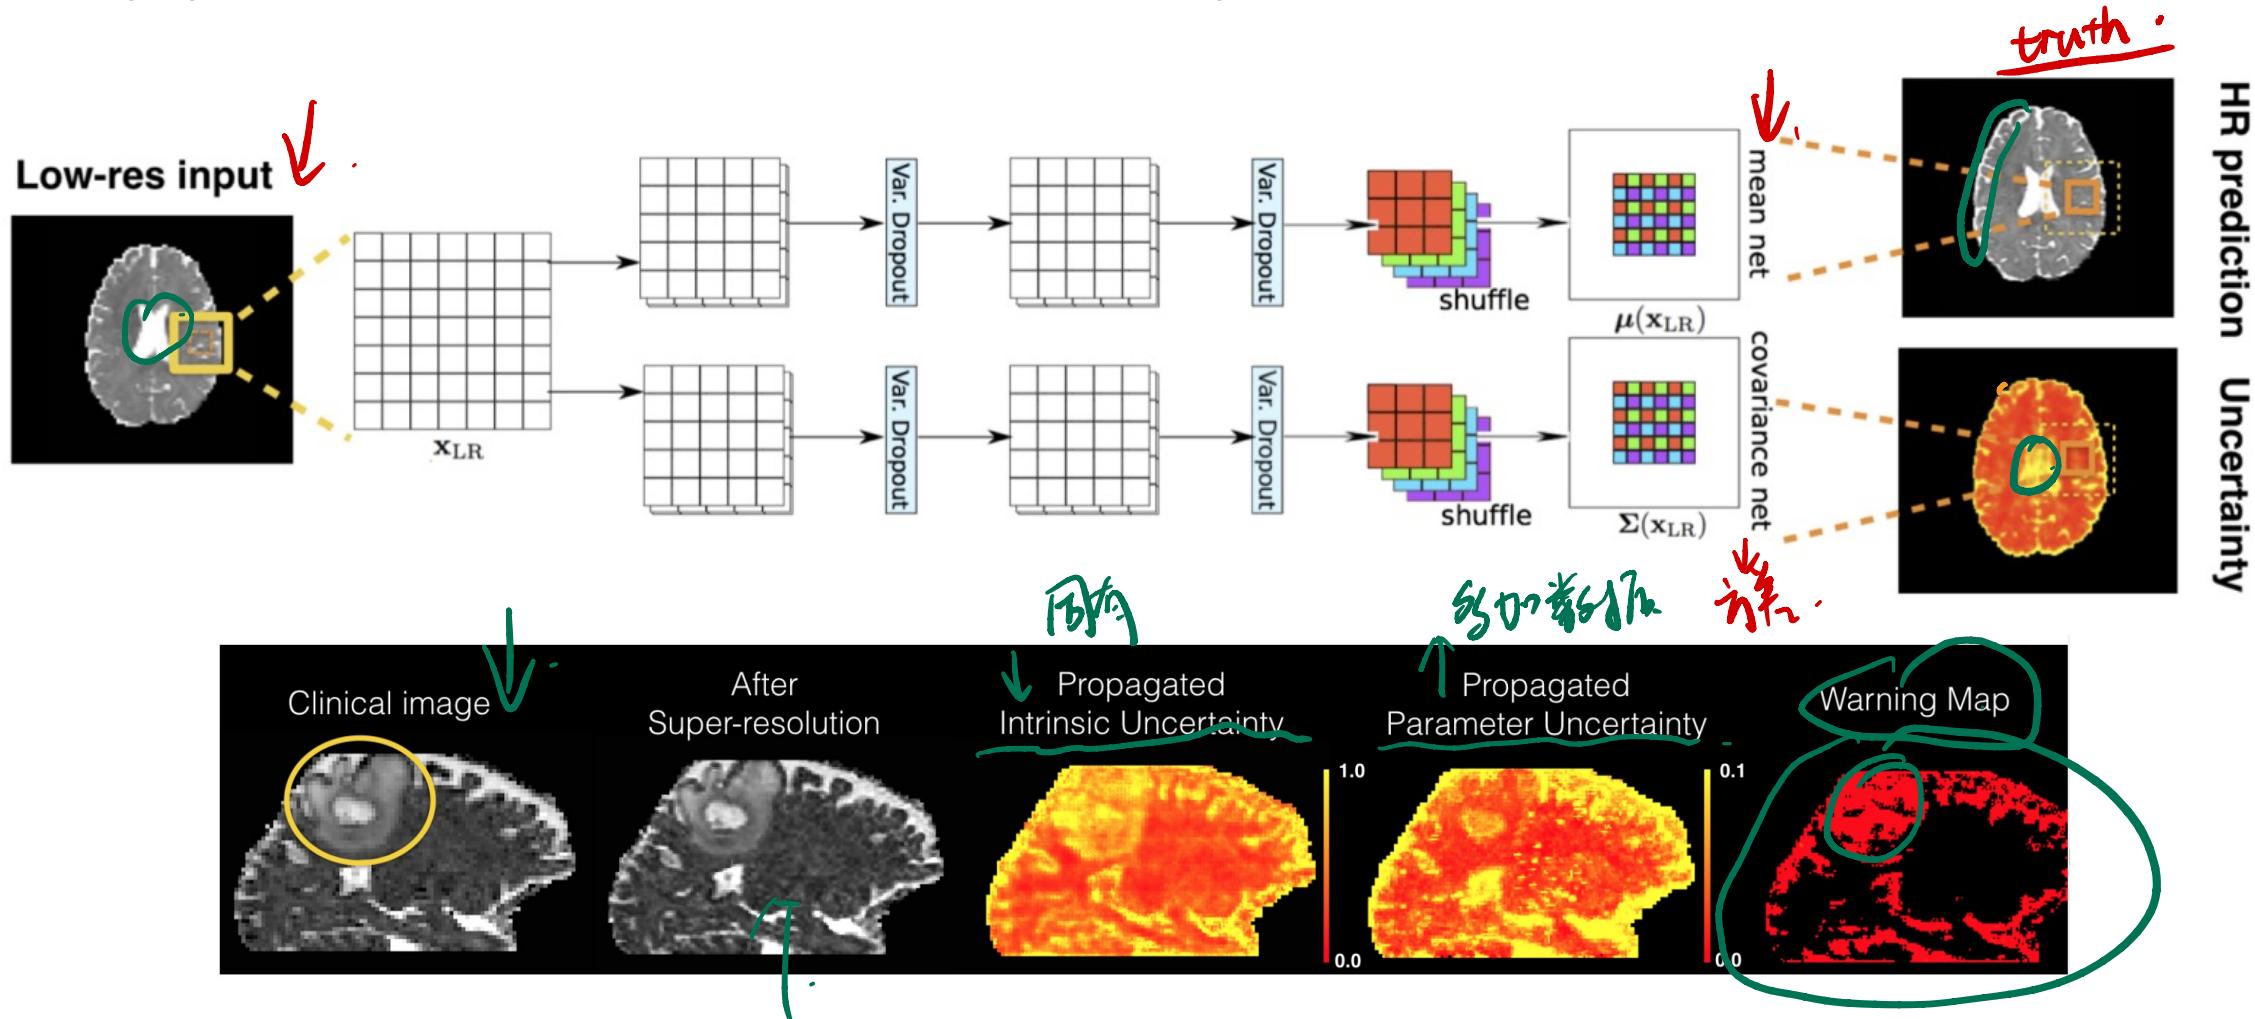
Part IV: Advances & Future Works

- Various applications
- Overview of recent progresses
- Future directions

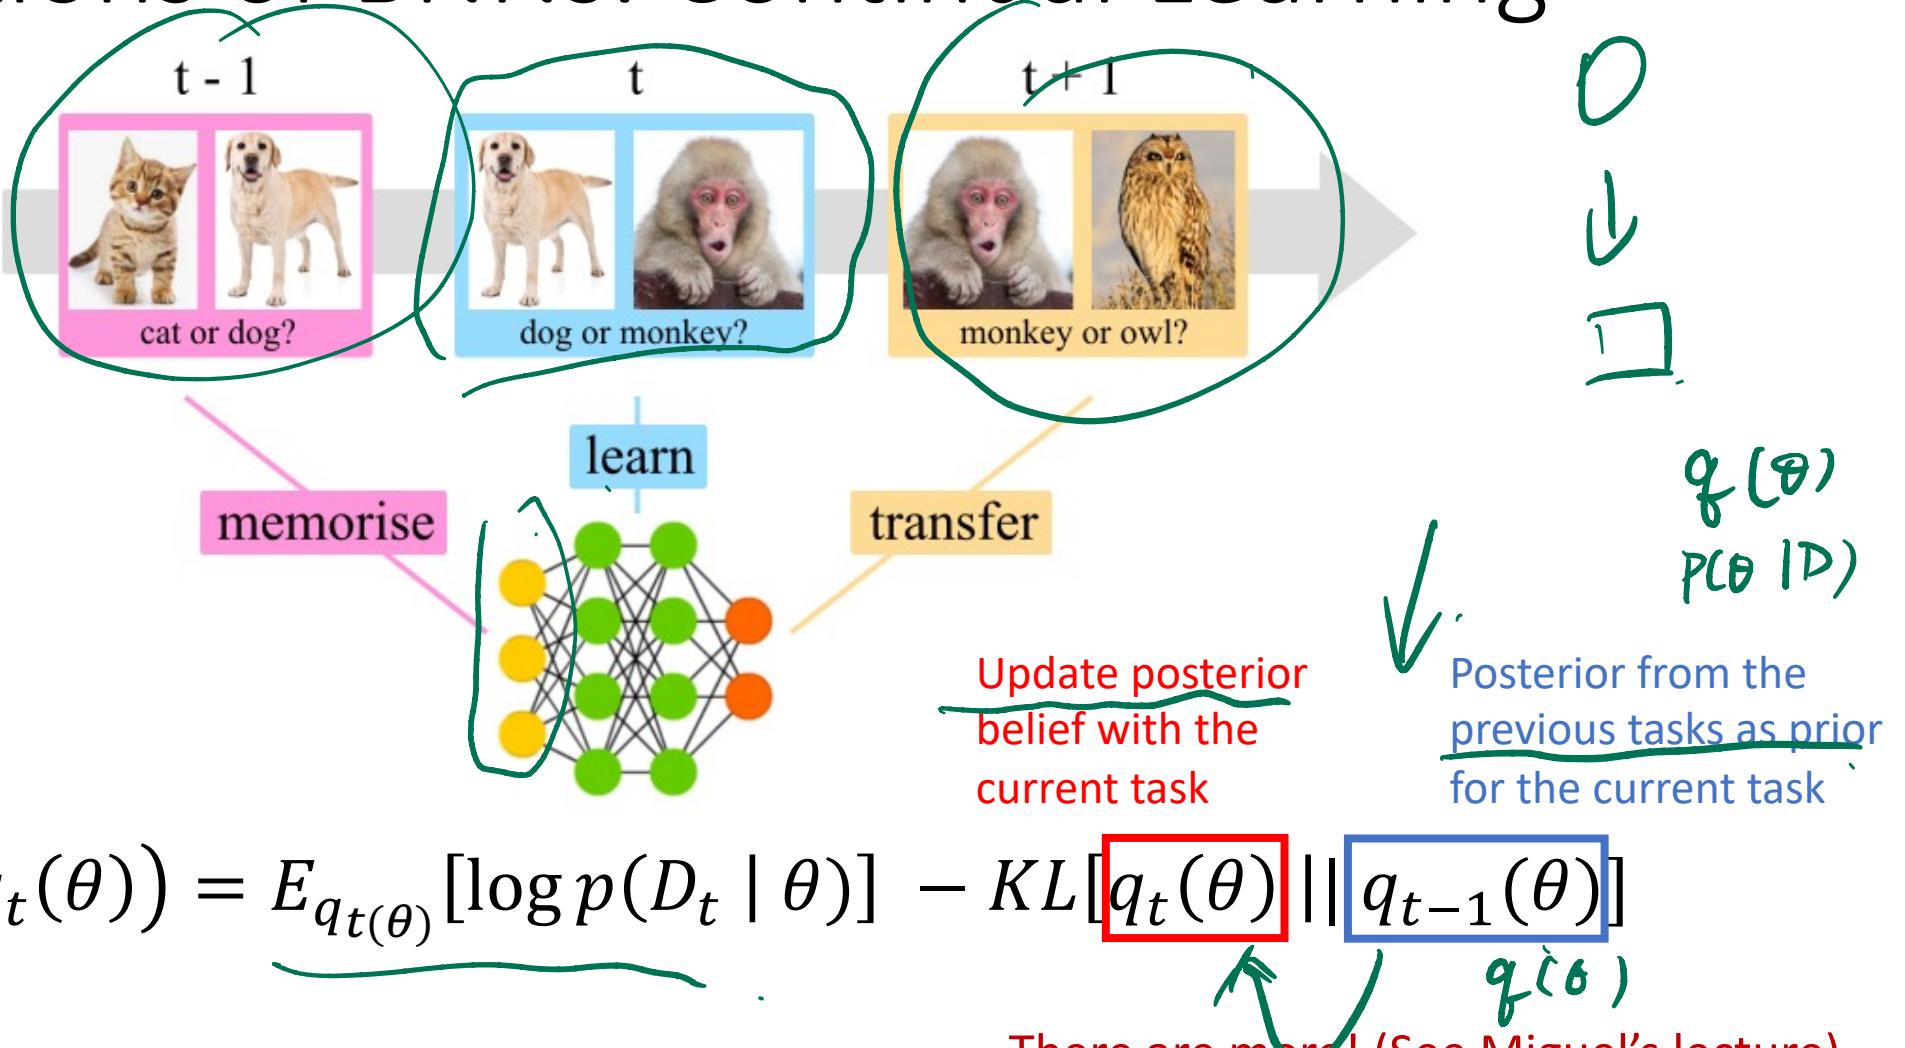
Applications of BNNs: Image Segmentation



Applications of BNNs: Super Resolution



Applications of BNNs: Continual Learning



Recent Progress in BNNs: Inference

$$\text{SGD: } \theta_{t+1} = \theta_t - \eta \nabla_{\theta_t} \tilde{U}(\theta_t)$$

$$\text{SGLD: } \theta_{t+1} = \theta_t - \eta \nabla_{\theta_t} \tilde{U}(\theta_t) + \sqrt{2\eta}\epsilon, \quad \epsilon \sim N(0, I)$$

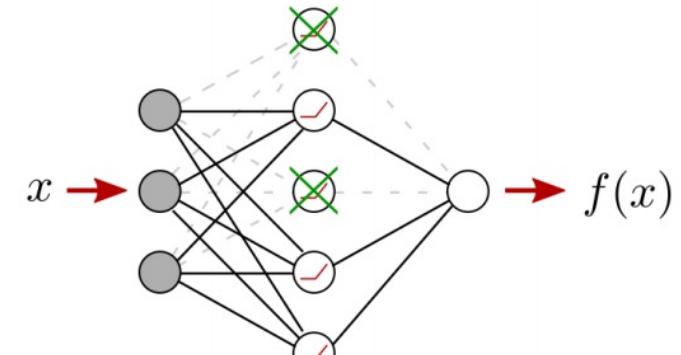
Stochastic gradient MCMC

Recent Progress in BNNs: Inference

$$\text{SGD: } \theta_{t+1} = \theta_t - \eta \nabla_{\theta_t} \tilde{U}(\theta_t)$$

$$\text{SGLD: } \theta_{t+1} = \theta_t - \eta \nabla_{\theta_t} \tilde{U}(\theta_t) + \sqrt{2\eta}\epsilon, \quad \epsilon \sim N(0, I)$$

Stochastic gradient MCMC

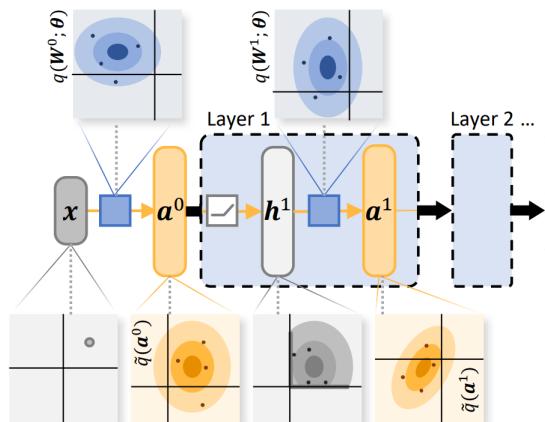


Recent Progress in BNNs: Inference

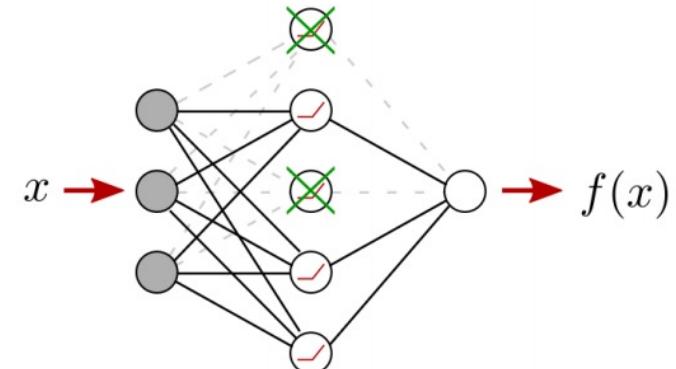
$$\text{SGD: } \theta_{t+1} = \theta_t - \eta \nabla_{\theta_t} \tilde{U}(\theta_t)$$

$$\text{SGLD: } \theta_{t+1} = \theta_t - \eta \nabla_{\theta_t} \tilde{U}(\theta_t) + \sqrt{2\eta}\epsilon, \quad \epsilon \sim N(0, I)$$

Stochastic gradient MCMC



Deterministic approximations



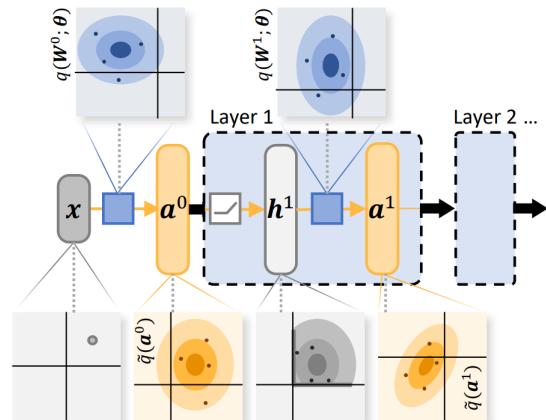
Monte Carlo dropout

Recent Progress in BNNs: Inference

$$\text{SGD: } \theta_{t+1} = \theta_t - \eta \nabla_{\theta_t} \tilde{U}(\theta_t)$$

$$\text{SGLD: } \theta_{t+1} = \theta_t - \eta \nabla_{\theta_t} \tilde{U}(\theta_t) + \sqrt{2\eta}\epsilon, \quad \epsilon \sim N(0, I)$$

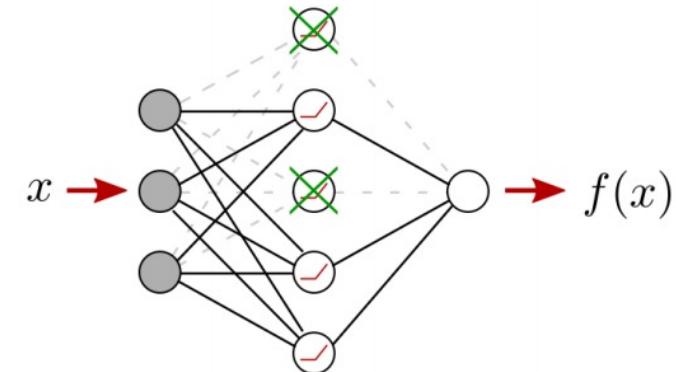
Stochastic gradient MCMC



Deterministic approximations

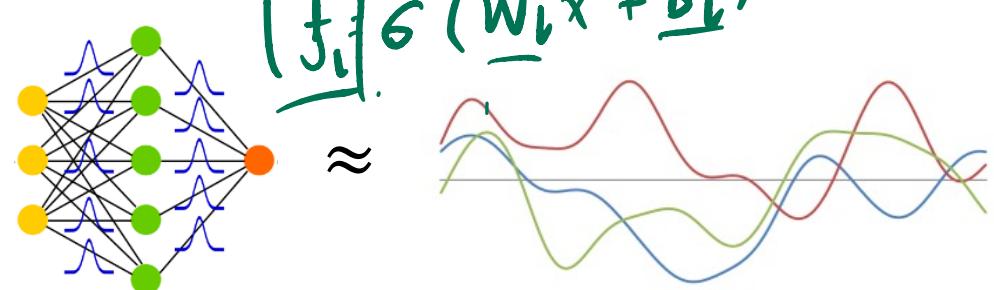
VI

① ~~dark~~
② ~~stable~~



Monte Carlo dropout

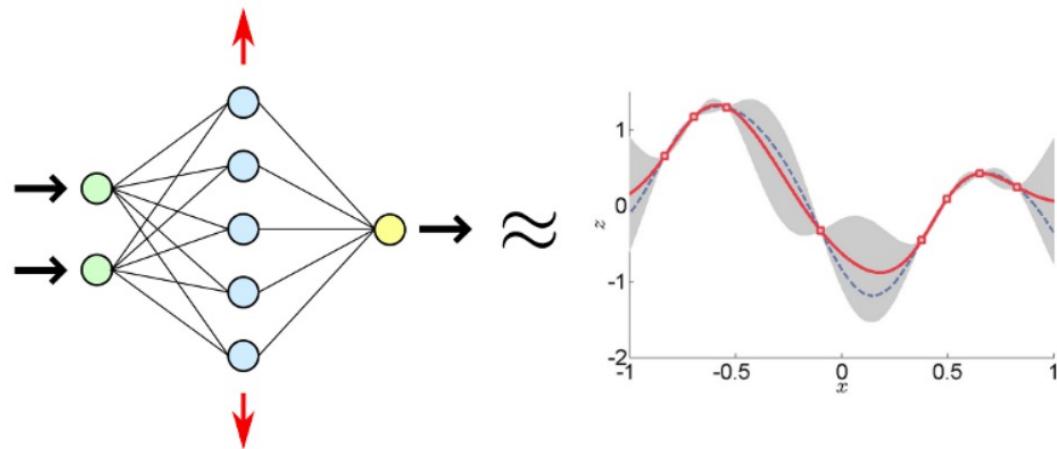
$f_b \sim \text{prior}$.



Function space approximate inference

There are more! (See Miguel's lecture)

Recent Progress in BNNs: Theory



Connections to GPs:

- BNN with very wide hidden layers ≈ Gaussian process
- Width limit convergence: in both prior (Neal's result) and posterior

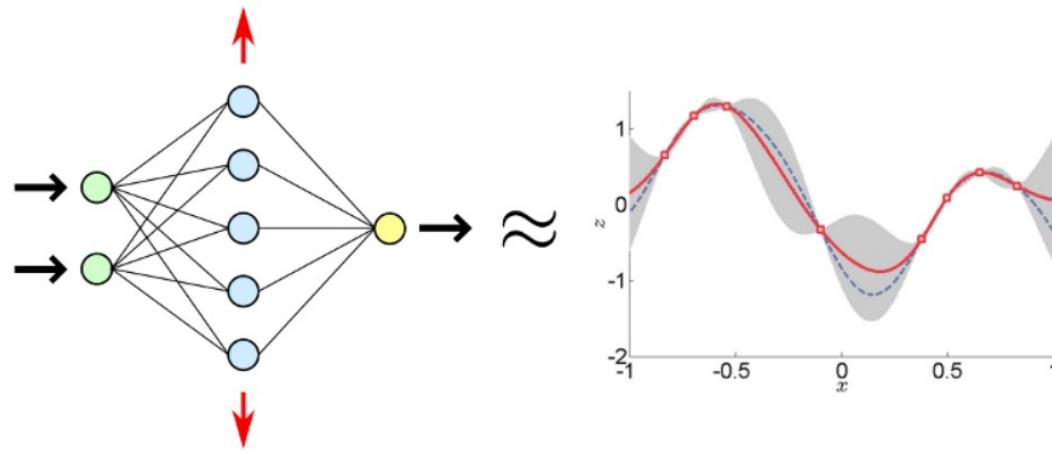
Neal. Bayesian Learning for Neural Networks. PhD Thesis, 1996

Matthews et al. Gaussian Process Behaviour in Wide Deep Neural Networks. ICLR 2018

Lee et al. Deep Neural Networks as Gaussian Processes. ICLR 2018

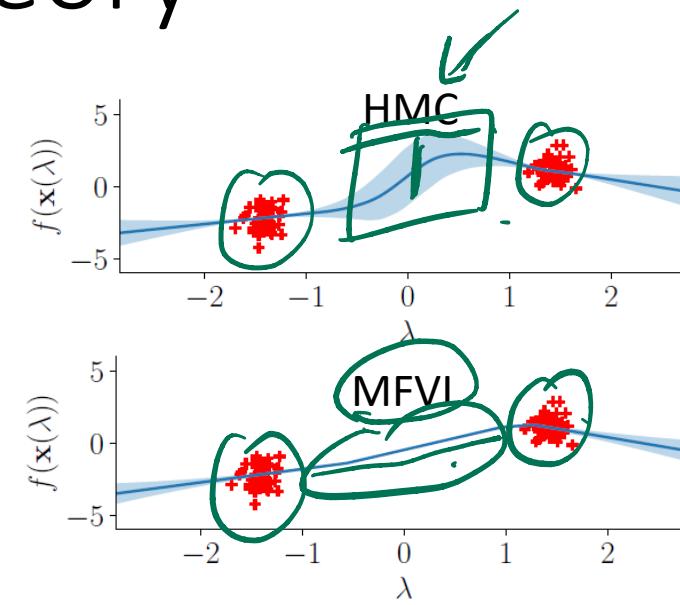
Hron et al. Exact posterior distributions of wide Bayesian neural networks. 2020

Recent Progress in BNNs: Theory



Connections to GPs:

- BNN with very wide hidden layers
≈ Gaussian process
- Width limit convergence: in both prior (Neal's result) and posterior



Approx. vs exact inference: underfit.

- Theoretical limitation of MFVI in shallow BNNs with ReLU activations
- Empirically deep BNNs with MFVI still fails in certain cases

Foong et al. On the Expressiveness of Approximate Inference in Bayesian Neural Networks. NeurIPS 2020

Farquhar et al. Liberty or Depth: Deep Bayesian Neural Nets Do Not Need Complex Weight Posterior Approximations. NeurIPS 2020

Coker et al. Wide Mean-Field Bayesian Neural Networks Ignore the Data. AISTATS 2022

Future directions

- Understanding BNN behaviour:
 - How would $q(f)$ behave given a particular form of $q(W)$? *function space*
 - Is weight-space objective appropriate for MFVI? *$g(w)$, MFVI*
 - We don't understand very well the optimisation properties of VI-BNN *VI*
- Computational complexity overhead: worth it?
 - How can we make the approximate posterior more efficient *#* *eff.*
 - in both time and space complexities?
- Priors for BNNs
 - “Default” Gaussian prior $N(W; 0, \sigma^2 I)$: the right prior? *f, w.*
 - How to think about priors in function space?
- Applications
 - Improve for applications that require good uncertainty estimates

Thank You!

Questions? Ask NOW or email:
yingzhen.li@imperial.ac.uk

Example answers of the tutorial demos:
Regression: <https://bit.ly/39eZHit>
Classification: <https://bit.ly/3QikcLO>